

COMP9418: Advanced Topics in Statistical Machine Learning

Jointree Algorithm

Instructor: Gustavo Batista


University of New South Wales

Introduction

- In this lecture, we will study a variation of VE known as *jointree algorithm*
 - Also known as *clique-tree* and *tree-clustering* algorithm
 - Jointree can be understood in terms of factor elimination
 - It improves VE complexity by answering multiple queries
 - It forms the basis of a class of approximate algorithms we will discuss later in this course
- We will start describing the idea of inference by factor elimination
 - In the sequence, we formalise the jointree algorithm using these ideas

Introduction

- Given a network we want to compute posterior marginals for each of its n variables
 - VE can compute a single marginal in $O(n \exp(w))$, where w is the width of the elimination order
 - We can run VE n times, leading to a total complexity of $O(n^2 \exp(w))$
 - The n^2 term can be problematic even when the treewidth is small
- Jointree can avoid this complexity leading to a $O(n \exp(w))$ time and space complexity
 - Bayesian networks, it will compute the posterior marginals for all network families (a variable and its parents)
 - For Markov networks, it will provide posterior marginals for all cliques (clique-tree algorithm)



Factor Elimination

- We want to compute prior marginals over some variable Q
 - VE eliminates every other variable
 - Factor elimination will eliminate all factors except for one that contain Q
- The elimination of factor f_i from a set of factors \mathcal{S} is a two-step process
 - Eliminate all variables V that appear *only* in factor f_i
 - Multiply the result $\sum_V f_i$ by some other factor f_j in the set \mathcal{S}

Factor Elimination Algorithm: FE1

Input: Network N , a variable Q in the network

Output: prior marginal $P(Q)$

$S \leftarrow$ factors of network N

$f_r \leftarrow$ a factor in S that contains variable Q

while S has more than one factor **do**

 remove a factor $f_i \neq f_r$ from set S

$V \leftarrow$ variables that appear in factor f_i but not in S

$f_j \leftarrow f_j \sum_V f_i$ for some factor $f_j \in S$

return $\text{project}(f_r, Q)$

- This algorithm makes use of the factor operation $\text{project}(f, Q)$, which simply sums out all variables not in Q :


$$\text{project}(f, Q) \stackrel{\text{def}}{=} \sum_{\text{vars}(f) - Q} f$$

Factor Elimination: Correctness

- Factor elimination is simply a variation of VE
 - While VE eliminates one variable at a time, FE eliminates a set of variables V at once
 - As these variables appear only in f_i , we replace f_i by a new factor $\sum_V f_i$
 - We take an extra step and multiply the new factor by other factor f_j
- At each iteration, the number of factors in S decreases by 1
 - After enough iterations, S will contain a single factor that contains Q
 - Eliminating all other variables but Q provides the answer to the query
- Two open choices
 - Which factor f_i to eliminate next
 - Which factor f_j to multiply by

} Any set of choices will provide a valid answer
But some choices will be computationally better

Factor Elimination: Example




Suppose we want to answer $P(C)$ for this Bayesian network

A	$f_A = \Theta(A)$
a	.6
\bar{a}	.4

A	B	$f_B = \Theta_{B A}$
a	b	.9
a	\bar{b}	.1
\bar{a}	b	.2
\bar{a}	\bar{b}	.8

B	C	$f_C = \Theta_{C B}$
b	c	.3
b	\bar{c}	.7
\bar{b}	c	.5
\bar{b}	\bar{c}	.5


Factor Elimination: Example



We start eliminating f_A .
As A is in f_A and f_B , $V = \emptyset$

A	$f_A = \Theta(A)$	\times	A	B	$f_B = \Theta_{B A}$		B	C	$f_C = \Theta_{C B}$
a	.6		a	b	.9		b	c	.3
\bar{a}	.4		a	\bar{b}	.1		b	\bar{c}	.7
			\bar{a}	b	.2		\bar{b}	c	.5
			\bar{a}	\bar{b}	.8		\bar{b}	\bar{c}	.5

Factor Elimination: Example




We start eliminating f_A .
As A is in f_A and f_B , $V = \emptyset$

A	$f_A = \Theta(A)$	\times	A	B	$f_B = \Theta_{B A}$		B	C	$f_C = \Theta_{C B}$
a	.6		a	b	.9		b	c	.3
\bar{a}	.4		a	\bar{b}	.1		b	\bar{c}	.7
			\bar{a}	b	.2		\bar{b}	c	.5
			\bar{a}	\bar{b}	.8		\bar{b}	\bar{c}	.5

A	B	$f_A f_B$		B	C	f_C
a	b	.54		b	c	.3
a	\bar{b}	.06		b	\bar{c}	.7
\bar{a}	b	.08		\bar{b}	c	.5
\bar{a}	\bar{b}	.32		\bar{b}	\bar{c}	.5

Factor Elimination: Example



Now, we eliminate $f_A f_B$.
 $V = \{A\}$

A	B	$f_A f_B$
a	b	.54
a	\bar{b}	.06
\bar{a}	b	.08
\bar{a}	\bar{b}	.32


B	C	f_C
b	c	.3
b	\bar{c}	.7
\bar{b}	c	.5
\bar{b}	\bar{c}	.5

B	$\sum_A f_A f_B$
b	.62
\bar{b}	.38

X

B	C	f_C
b	c	.3
b	\bar{c}	.7
\bar{b}	c	.5
\bar{b}	\bar{c}	.5

Factor Elimination: Example




We multiply $\sum_A f_A f_B$ into f_C

B	$\sum_A f_A f_B$		B	C	f_C
b	.62		b	c	.3
\bar{b}	.38	X	b	\bar{c}	.7
			\bar{b}	c	.5
			\bar{b}	\bar{c}	.5

B	C	$f_C \sum_A f_A f_B$
b	c	.186
b	\bar{c}	.434
\bar{b}	c	.190
\bar{b}	\bar{c}	.190

Factor Elimination: Example





Finally, we eliminate all other variables to obtain the answer

B	$\sum_A f_A f_B$		B	C	f_C
b	.62		b	c	.3
\bar{b}	.38	X	b	\bar{c}	.7
			\bar{b}	c	.5
			\bar{b}	\bar{c}	.5

B	C	$f_C \sum_A f_A f_B$		C	$\sum_B f_C \sum_A f_A f_B$
b	c	.186		c	.376
b	\bar{c}	.434	→	\bar{c}	.624
\bar{b}	c	.190			
\bar{b}	\bar{c}	.190			



Elimination Trees

- In variable elimination, the elimination order was used to specify an elimination strategy
 - The amount of work performed by VE was determined by the quality (width) of the order
- In factor elimination, we use trees to specify an elimination strategy
 - Each organization of factors into a tree structure represents a particular strategy
 - The quality of such trees (also called *width*) can be used to quantify the amount of work performed
- The figure shows one such tree
 - We call it *elimination tree*




Elimination Trees

- An elimination tree for a set of factors \mathcal{S} is a pair (T, ϕ) where T is a tree
 - Each factor in \mathcal{S} is assigned to exactly one node in T
 - We use ϕ_i to denote the product of factors assigned to node i in T
 - We also use $\text{vars}(i)$ to denote the variables in factor ϕ_i
- In this figure, the elimination tree (T, ϕ) has five nodes which are in one-to-one correspondence with the given factors
 - A node in an elimination tree may have multiple factors assigned to it or no factors at all
 - For many examples, there will be a one-to-one correspondence between factors and nodes in an elimination tree





Elimination Trees

- When computing the marginal over variables Q , we need to choose a special node r
 - This node r , called a *root*, must be chosen such that $Q \subseteq \text{vars}(r)$
 - For example, if $Q = \{C\}$, we can choose nodes 3, 4 and 5 can act as roots
 - A root node is not strictly needed but will simplify the discussion



Elimination Trees

- Given an elimination tree and a corresponding root r , our elimination strategy is
 - Eliminate factor ϕ_i only if it has a single neighbor j and $i \neq r$
 - Sum out variables V that appear in ϕ_i but not in the rest of the tree
 - Multiply the result $\sum_V \phi_i$ into the factor ϕ_j associated with its single neighbor j



Factor Elimination Algorithm: FE2

Input: Network N , a set of variables Q in the network, an elimination tree (T, ϕ) , a root node r

Output: prior marginal $P(Q)$

while tree T has more than one node **do**

 remove a node $i \neq r$ having a single neighbour j from T


$V \leftarrow$ variables appearing in ϕ_i but not in remaining tree T

$\phi_j \leftarrow \phi_j \sum_V \phi_i$


return project(ϕ_r, Q)

- We still need to make a choice of which node to remove since we may have more than one node i in the tree that satisfies the stated properties
- However, the choice made at this step does not affect the amount of work done by the algorithm


Elimination Trees: Example



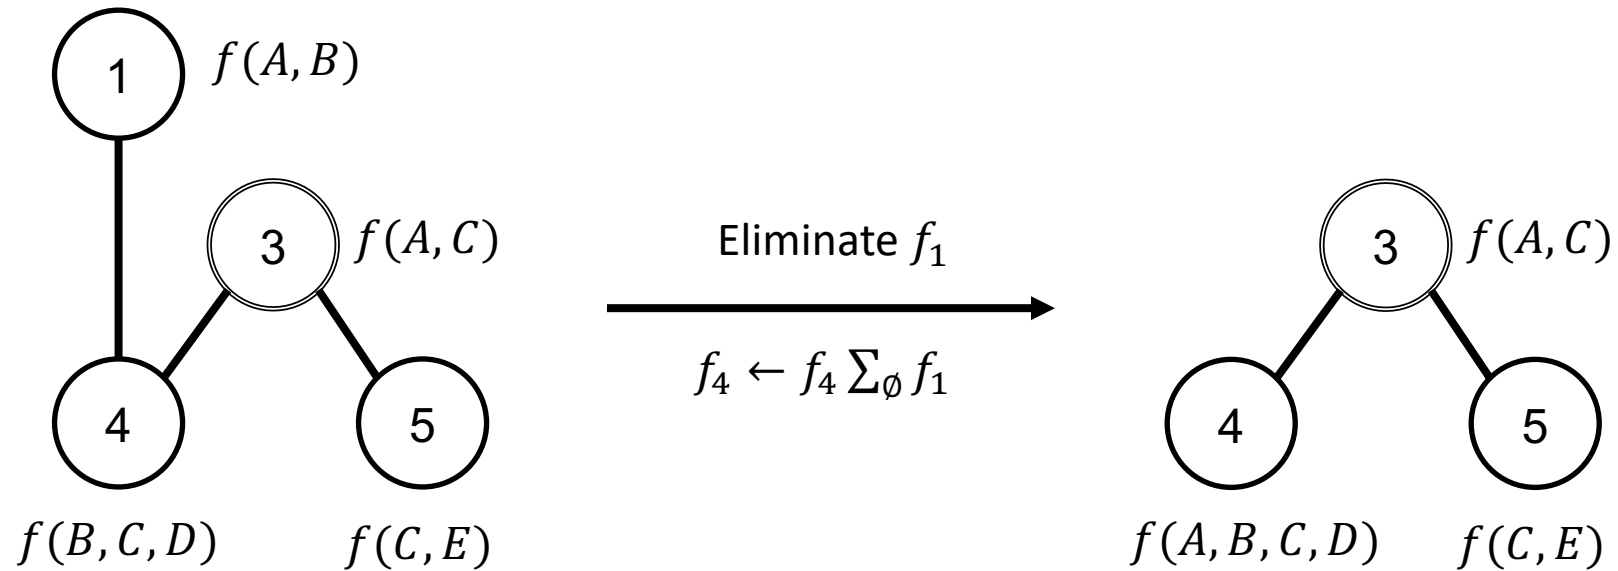

Let us suppose we want to compute the marginal over variable C . We set $r = 3$



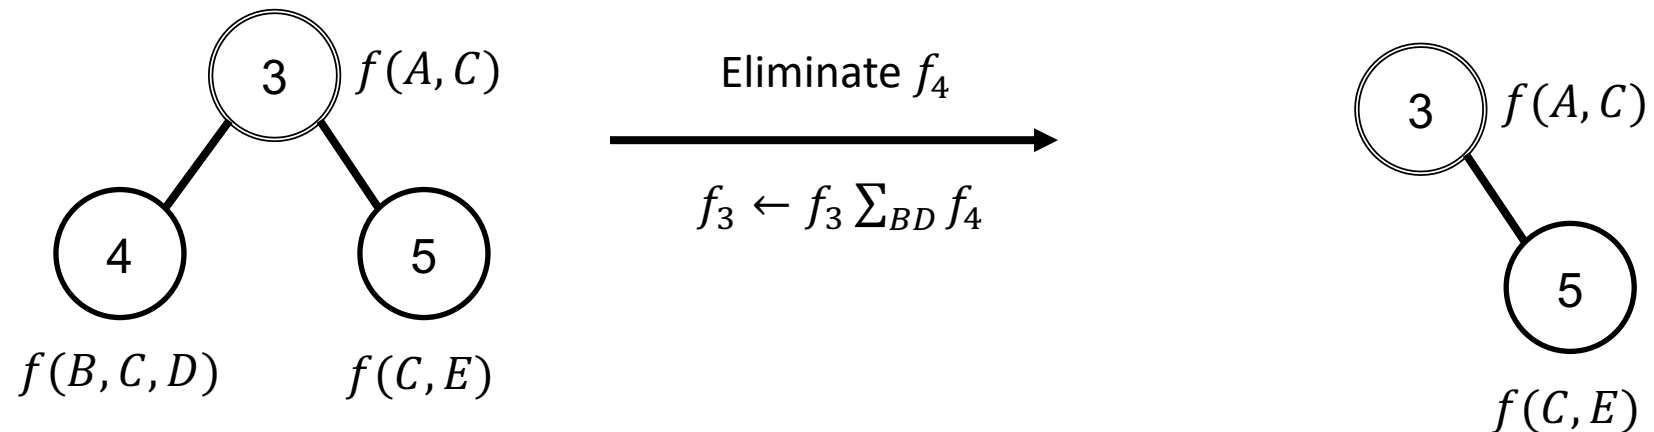

Eliminate f_2
 $f_1 \leftarrow f_1 \Sigma_{\emptyset} f_2$





Elimination Trees: Example




Elimination Trees: Example



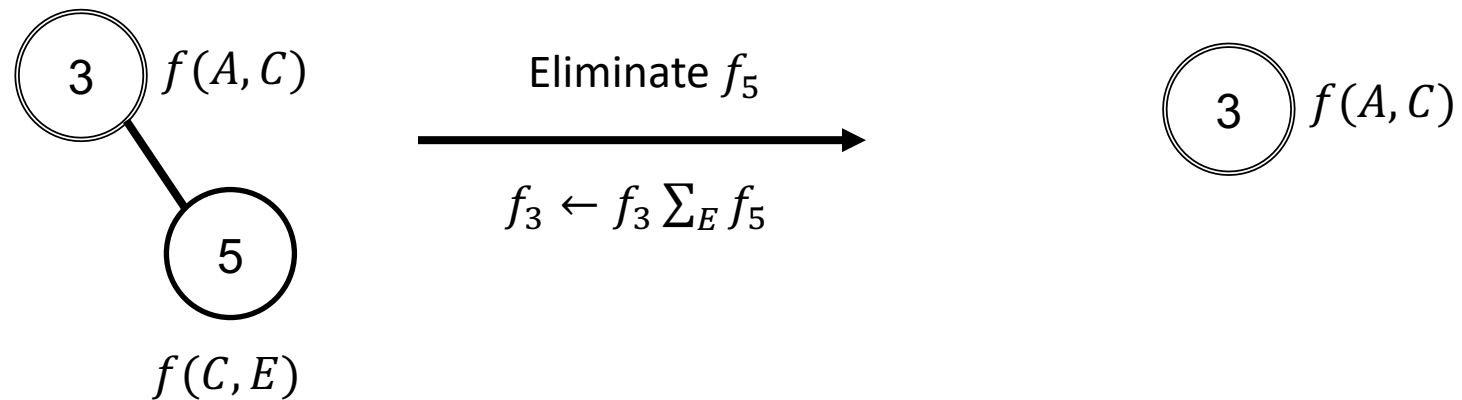
Elimination Trees: Example



Elimination Trees: Example





The final factor is over variables A, C .
Eliminating A gives the desired result.



Elimination Trees and Runtime

- With VE, any elimination order will lead to correct results
 - Yet a specific order may have a considerable effect on the amount of work
- FE is similar
 - Any elimination tree will lead to correct results
 - Yet some trees will lead to less work
 - We will return to this topic later




Separator



- The separator of edge $i - j$ in an elimination tree is a set of variables defined as follows

$$\mathcal{S}_{ij} = vars(i, j) \cap vars(j, i)$$

- $vars(i, j)$ are variables that appear on the i -side of edge $i - j$
 - $vars(j, i)$ are variables that appear on the j -side of edge $i - j$
- When variables V are summed out of factor f_i before it is eliminated, the resulting factor is guaranteed to be over separator \mathcal{S}_{ij}




Separator: Example

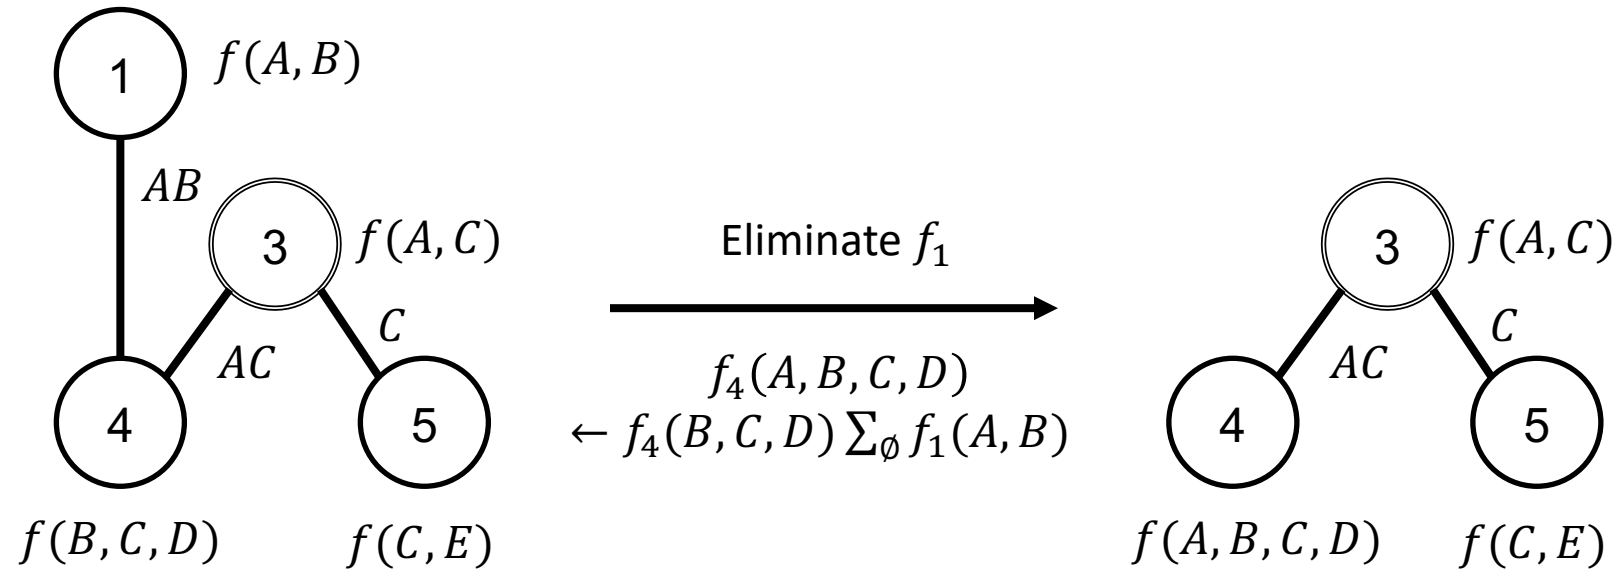



Eliminate f_2



$$f_1(A, B) \leftarrow f_1(A) \sum_{\emptyset} f_2(A, B)$$



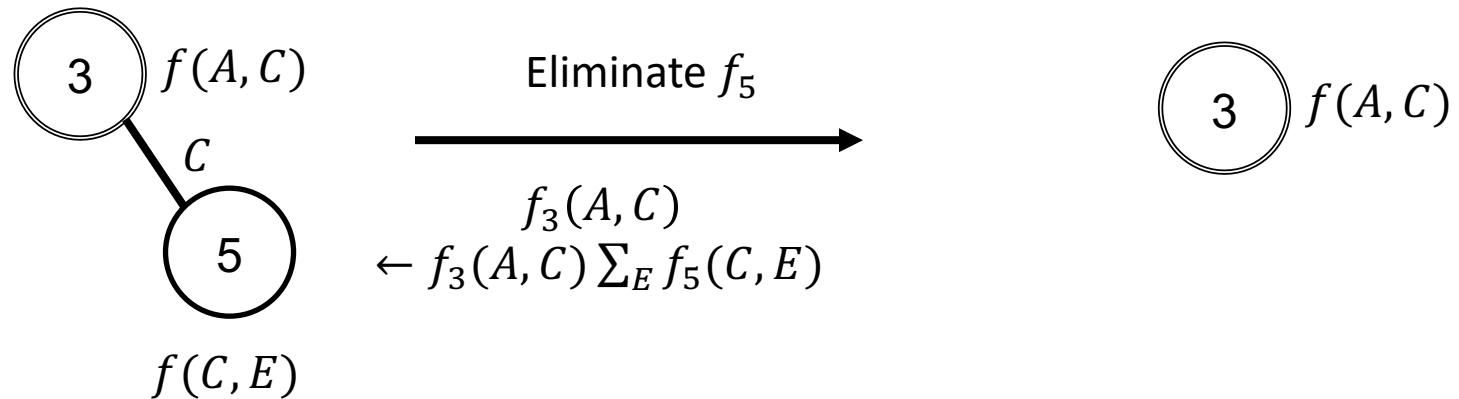

Separator: Example



Separator: Example



Separator: Example





Cluster

- The cluster of a node i in an elimination tree is a set of variables defined as follows:

$$C_i = vars(i) \cup \bigcup_j s_{ij}$$

- The width of an elimination tree is the size of its largest cluster – 1
- This elimination tree has width = 3





Cluster

- The cluster of a node i in an elimination tree is a set of variables defined as follows:



$$C_i = vars(i) \cup \bigcup_j s_{ij}$$

- The width of an elimination tree is the size of its largest cluster – 1
- This elimination tree has width = 3
- And this one has width = 2



Cluster

- Two key observations about clusters
 - When we are about to eliminate node i , the variables of factor ϕ_i are exactly the cluster of node i , C_i .
 - The factor ϕ_r must be over the cluster of root r , C_r
- Hence, FE2 can be used to compute the marginal over any subset of cluster C_r
 - These observations allow us to rephrase FE2
 - The new formulation takes advantages of both separators and clusters



Factor Elimination Algorithm: FE3

Input: Network N , a set of variables Q in the network, an elimination tree (T, ϕ) , a root node r in T where $Q \subseteq C_r$,

Output: prior marginal $P(Q)$

C_i is the cluster of node i in tree T

S_{ij} is the separator of edge $i - j$ in tree T

while tree T has more than one node **do**

 remove a node $i \neq r$ having a single neighbour j from T

$\phi_j \leftarrow \phi_j \text{project}(\phi_i, S_{ij})$


return $\text{project}(\phi_r, Q)$

Message-passing Formulation

- We will now rephrase our algorithm using a message passing paradigm
 - It will allow us to execute the algorithm without destroying the elimination tree in the process
 - This is important when computing multiple marginals
 - We can save intermediate computations and reuse them across different queries
- This reuse will be the key to achieving the complexity
 - Given an elimination tree of width w we will be able to compute the marginal over every cluster in $O(m \exp(w))$ time and space, where m is the number of nodes in the elimination tree

Message-passing Formulation


- Given an elimination tree (T, ϕ) with root r
 - For each node $i \neq r$ in the elimination tree, there is a unique neighbor of i that is closest to root r
 - A node i will be eliminated from the tree only after all its neighbors, except the one closest to the root, have been eliminated
 - When a node i is about to be eliminated, it will have a single neighbor j . Its current factor will have all variables but the separator S_{ij} eliminated and it will be multiplied by factor j



Elimination tree with directed edges pointing to neighbour closest to root node 7

Message-passing Formulation


- Now, we view the elimination of node i with a single neighbor j as a process of passing a message M_{ij} from node i to j
 - When j receives a message, it multiplies it into its current factor ϕ_j
 - Node i cannot send a message to j until it has received all messages from neighbors $k \neq j$
 - After i receives these messages, its current factor will be $\phi_i \prod_{k \neq j} M_{ki}$
 - The message i send to j will be $\sum_{C_i \setminus S_{ij}} \phi_i \prod_{k \neq j} M_{ki}$



Elimination tree with directed edges pointing to neighbour closest to root node 7

Message-passing Algorithm


- We can now formulate FE as message-passing algorithm
 - To compute the marginal over some variables Q
 - Select a root r in the elimination tree such that $Q \subseteq \mathcal{C}_r$
 - Push messages towards the root r
 - When all messages into the root are available, we multiply them by ϕ_r and eliminate variables not in Q
 - If our elimination tree has m nodes and $m - 1$ edges, then $m - 1$ messages need to be sent



10 messages are sent toward root node 7 in this 11-node elimination tree


Message-passing and Reuse

- Suppose we want to compute the marginal over some other cluster $C_i, i \neq r$
 - We choose i as the new root and repeat the message-passing process
 - Some additional messages need to be passed, but not as many as $m - 1$, assuming we saved the messages sent to r
 - In the figure, out 10 messages sent toward node 6, eight messages have already been computed when 7 was the root



Message-passing and Reuse

- The key observation is that if we choose every node as a root, the total number of messages is $2(m - 1)$
 - There are $m - 1$ edges and two distinct messages per edge
 - These messages are usually computed in two phases
 - Inward: we direct messages toward a root r
 - Outward: we direct messages away from the root r



Outward phase with node 7 as root

Message-passing: Complexity

- A message from i to j is computed by multiplying a few factors and summing out the separator
 - The factor that results from the multiplication must be over the cluster of node i
 - Hence, the complexity of both multiplication and summation is $O(\exp(w))$, where w is the size of cluster C_i
 - The width of an elimination tree is the size of its maximal cluster -1
 - Hence, if w is the width, then the cost of any message is $O(\exp(w))$
 - Since we have $2(m - 1)$, the total cost is $O(m \exp(w))$

Joint Marginals and Evidence

- Given some evidence e we want to use factor elimination to compute the joint marginal $P(C_i, e)$ for each cluster C_i in the elimination tree, we can
 1. Reduce each factor f given the evidence e , leading to a set of reduced factors f^e
 2. Introduce an evidence indicator λ_E for every variable E in evidence e . λ_E is a factor over variable E that captures the value of E in evidence e : $\lambda_E(e) = 1$ if e is consistent with evidence e and $\lambda_E(e) = 0$ otherwise

$e: \{A = \text{true}, B = \text{false}\}$

A	λ_A
a	1
\bar{a}	0

B	λ_B
b	0
\bar{b}	1

Joint Marginals and Evidence

- The first method is more efficient if we plan to compute marginals with respect to only one piece of evidence e
- The second method is more efficient to compute marginals with respect to multiple pieces of evidence, e_1, \dots, e_n
 - While trying to reuse messages across different pieces of evidence
 - This method is implemented by assigning the evidence indicator λ_E to a node i in the elimination tree while ensuring that $E \in C_i$.
 - As a result, the clusters and separators of the elimination tree will remain intact and so will its width

$e: \{A = \text{true}, B = \text{false}\}$

A	λ_A
a	1
\bar{a}	0

B	λ_B
b	0
\bar{b}	1

Factor Elimination Algorithm: FE

Input: Network N , a set of variables Q in the network, an elimination tree (T, ϕ)

Output: joint marginal $P(C_i, e)$ for each node i in the elimination tree

for each variable E in evidence e **do**

$i \leftarrow$ node in tree T such that $E \in C_i$

$\lambda_E \leftarrow$ evidence indicator for variable E # $\lambda_E = 1$ if $e \sim e$ and $\lambda_E(e) = 0$ otherwise

$\phi_i \leftarrow \phi_i \lambda_E$ # entering evidence at node i

choose a root node r in the tree T

pull/collect messages towards root r


push/distribute messages away from root r

return $\phi_i \prod_k M_{ki}$ for each node i in the tree T # joint marginal $P(C_i, e)$

- This algorithm uses the second method for accommodating evidence
- It computes joint marginals using two phases of message passing
- If we save the messages across different runs of the algorithm, then we can reuse these messages as long as they are not invalidated when the evidence changes
- When the evidence at node i changes, we need to invalidate all messages that depend on the factor at that node
- These messages happen to be the ones directed away from node i in the elimination tree

Polytree Algorithm

- When the network has a polytree structure
 - We can use an elimination tree that corresponds to the polytree structure
 - This special case of the algorithm is known as polytree algorithm or belief propagation algorithm. We will discuss them later
 - If k is the maximum number of parents in any node in the polytree, then k is also the width of the elimination tree
 - The time and space complexity are $O(n \exp(k))$, where n is the number of nodes in the polytree

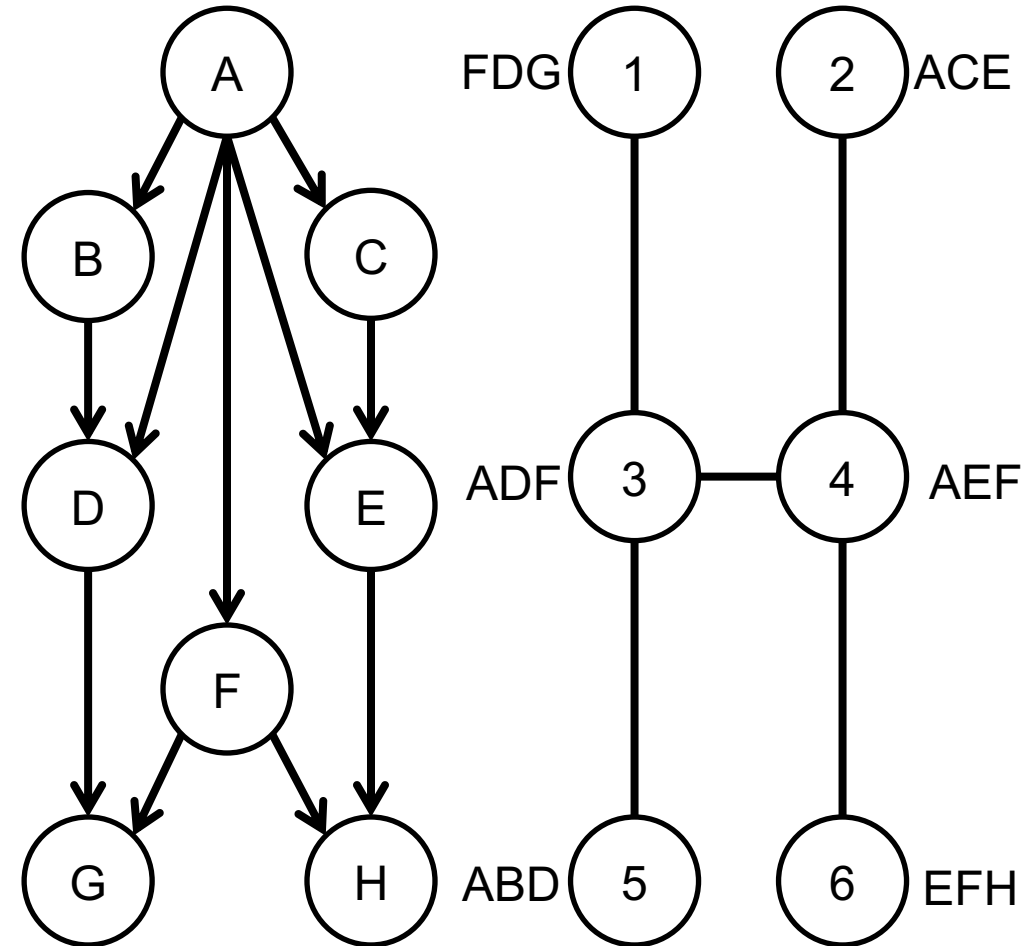


Jointree

- There are different methods for constructing elimination trees
 - But the method we discuss next will be based on an influential tool known as a *jointree*
 - It is this tool that gives factor elimination its traditional name: *the jointree algorithm*
- It is possible to phrase the factor elimination algorithm directly on jointrees without explicit mention elimination trees
 - This is indeed how the algorithm is classically described and we provide such a description
 - We start defining a jointree

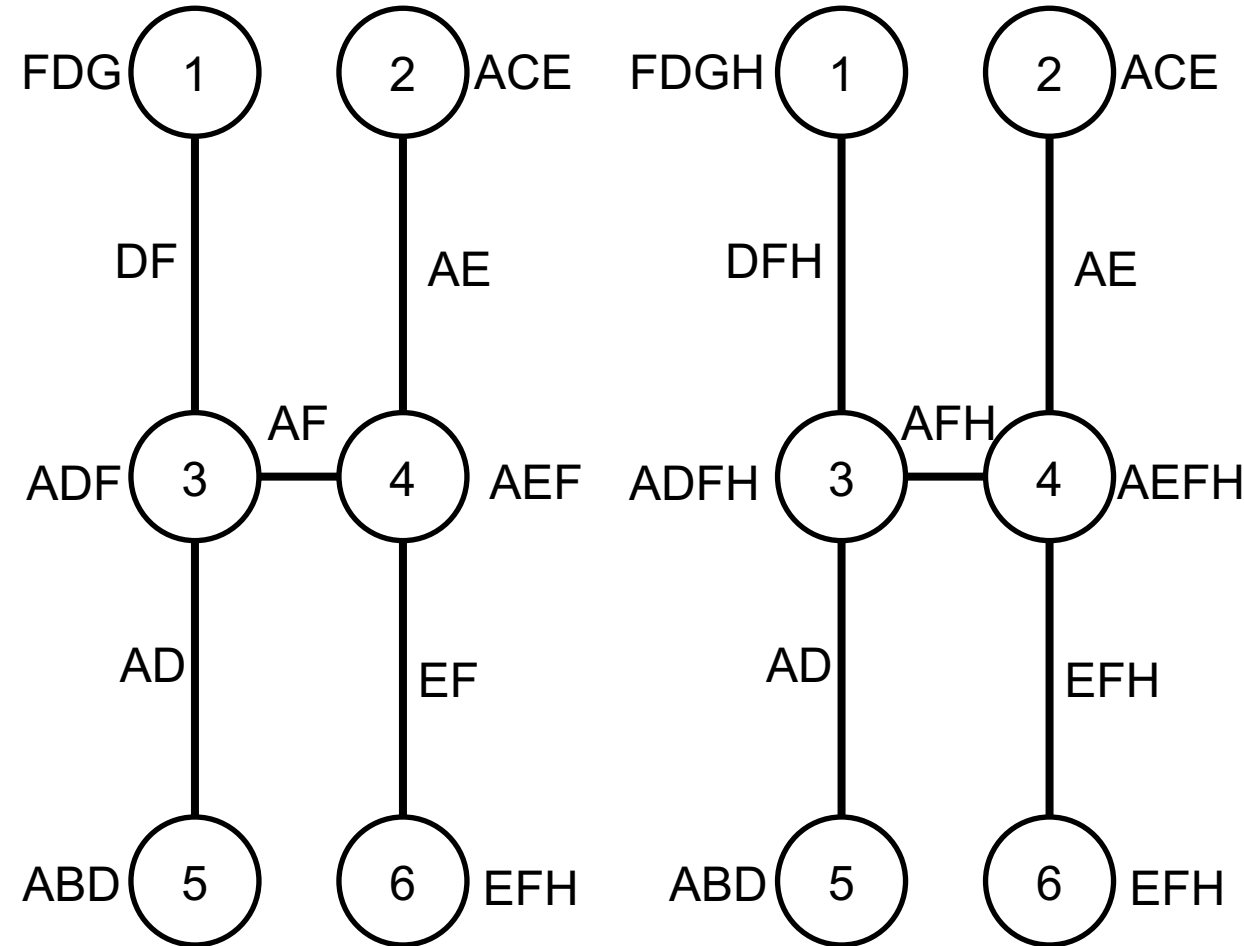
Jointree

- A *jointree* for a network G is a pair (T, \mathcal{C}) where T is a tree and \mathcal{C} is a function that maps each node i in the tree T into a label \mathcal{C}_i , called *cluster*.
- The jointree must satisfy the properties
 - The cluster \mathcal{C}_i is a set of nodes from G
 - Each factor in G must appear in some cluster \mathcal{C}_i
 - If a variable appears in two clusters \mathcal{C}_i and \mathcal{C}_j , it must appear in every cluster \mathcal{C}_k on the path connecting nodes i and j in the jointree. This is known as *jointree* or *running intersection* property



Jointree

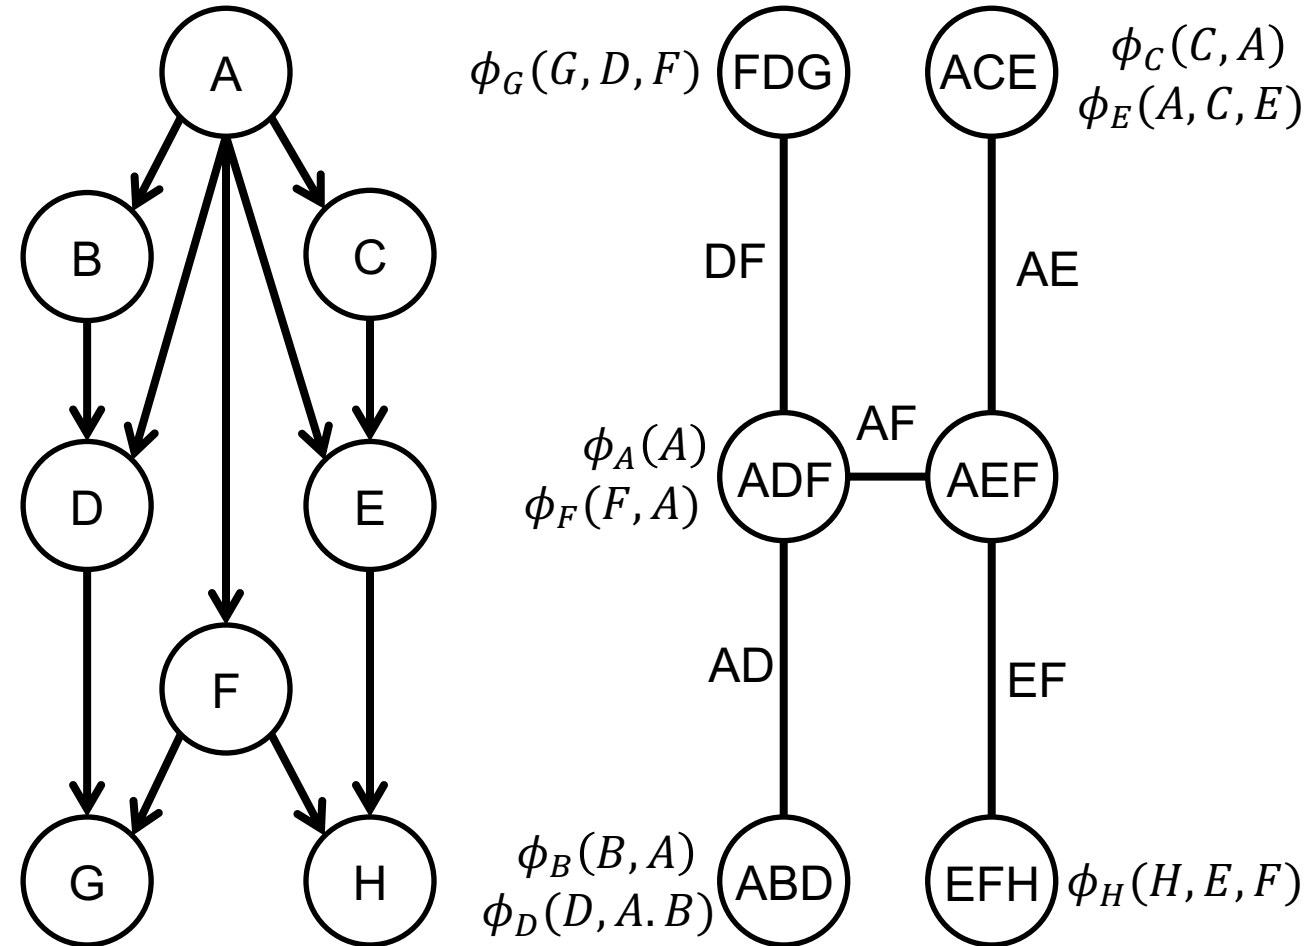
- The separator of edge $i - j$ in a jointree is denoted by S_{ij} and defined as $C_i \cap C_j$
 - The width of a jointree is defined as the size of its largest cluster minus one
- A jointree is *minimal* if it ceases to be a jointree for G once we remove a variable from one of its cluster
 - Left jointree is minimal but the right one is not



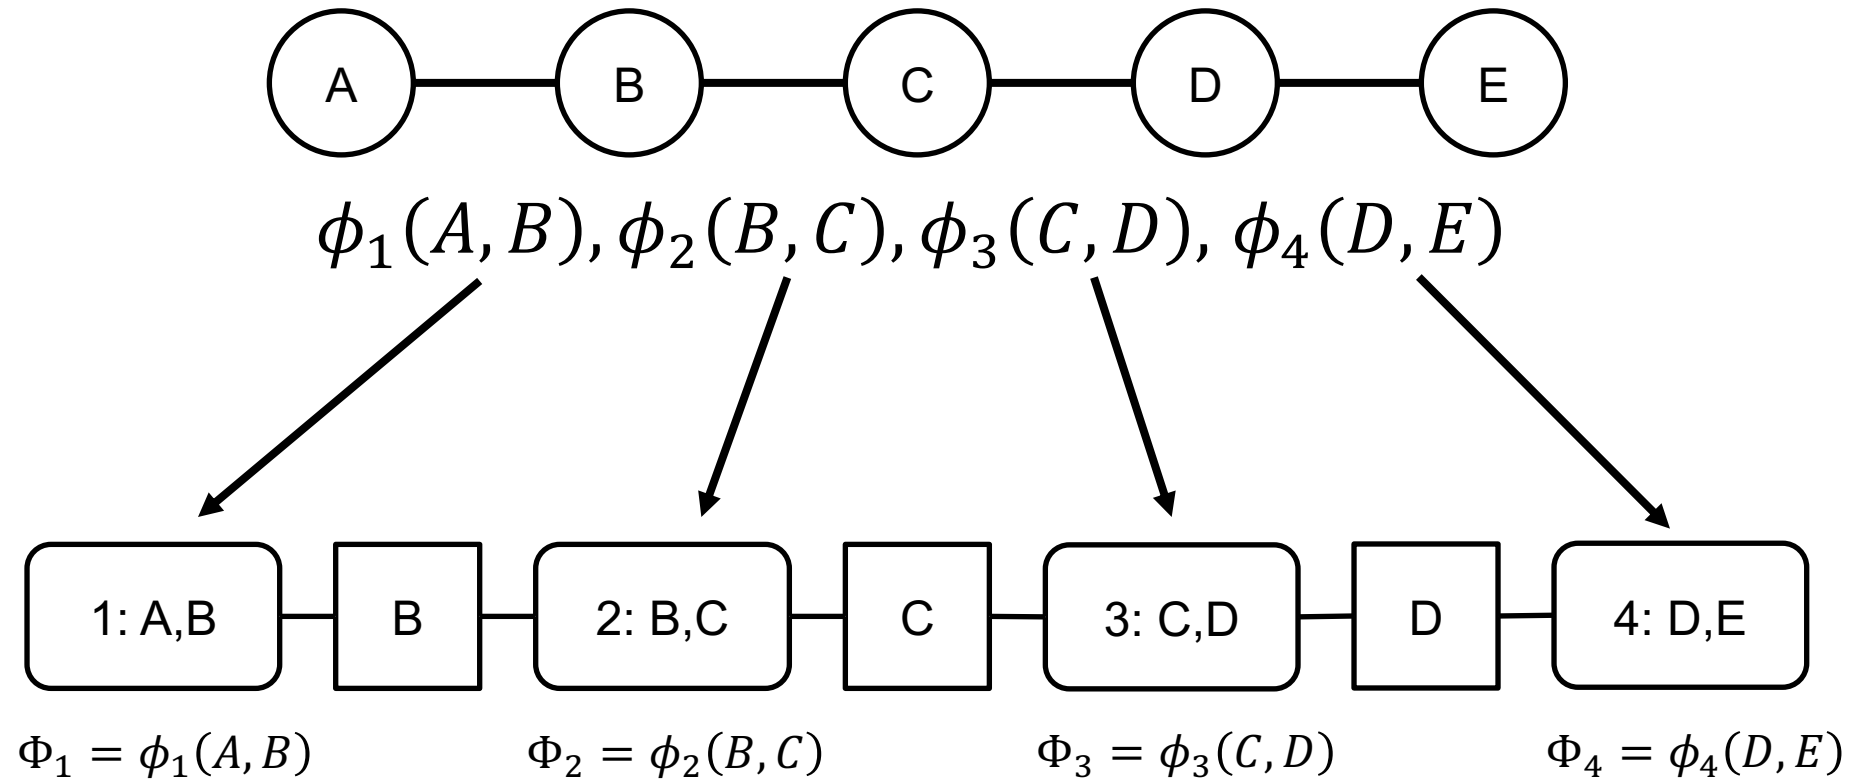
The Jointree Algorithm

- The classical jointree algorithm is:

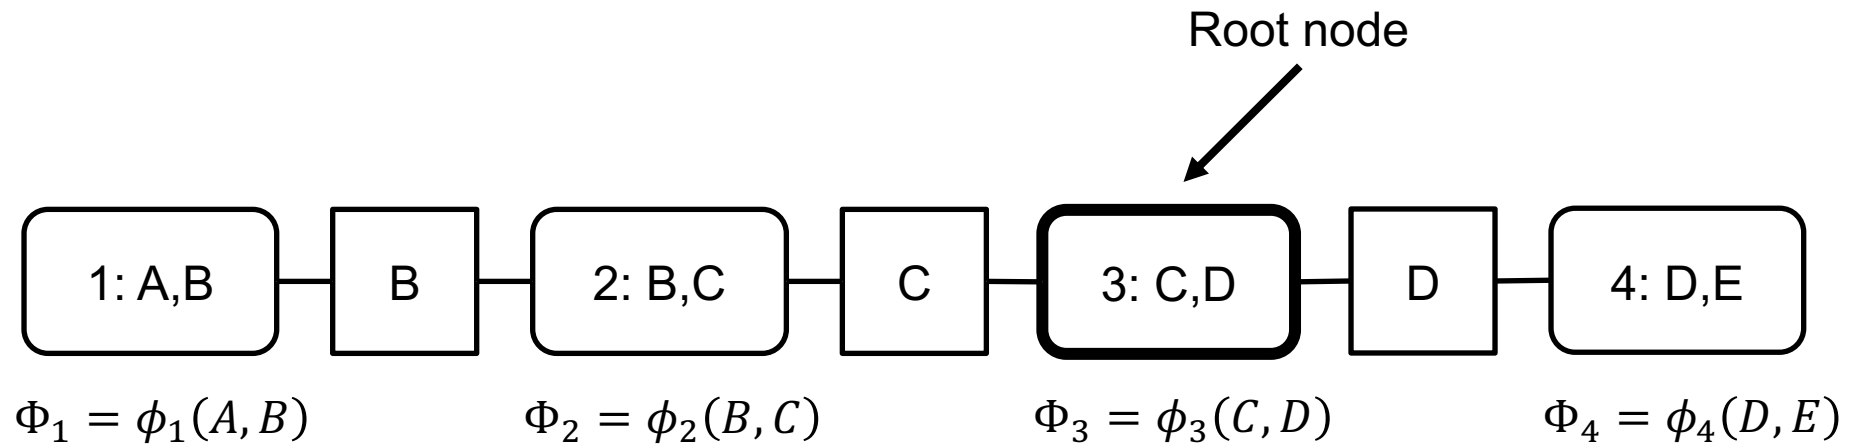

- Construct jointree (T, \mathcal{C}) for a given network
- Assign each factor ϕ_i to a cluster that contains $vars(\phi_i)$
- Assign each evidence indicator λ_X to a cluster that contains X
- Propagate messages in the jointree between the clusters
- After passing two messages per edge in the jointree, we can compute the marginals $P(\mathcal{C}, e)$ for every cluster \mathcal{C}





Message-passing: Example





Message-passing: Example



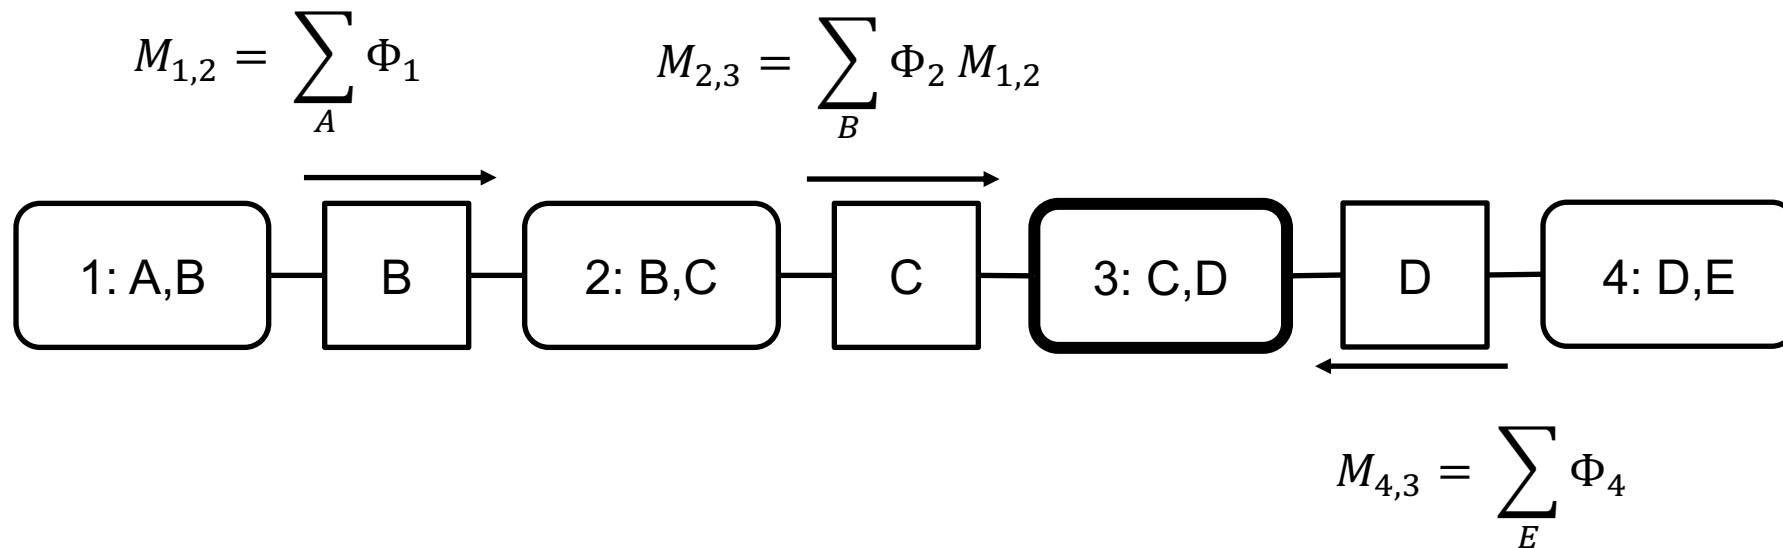

Message-passing: Example



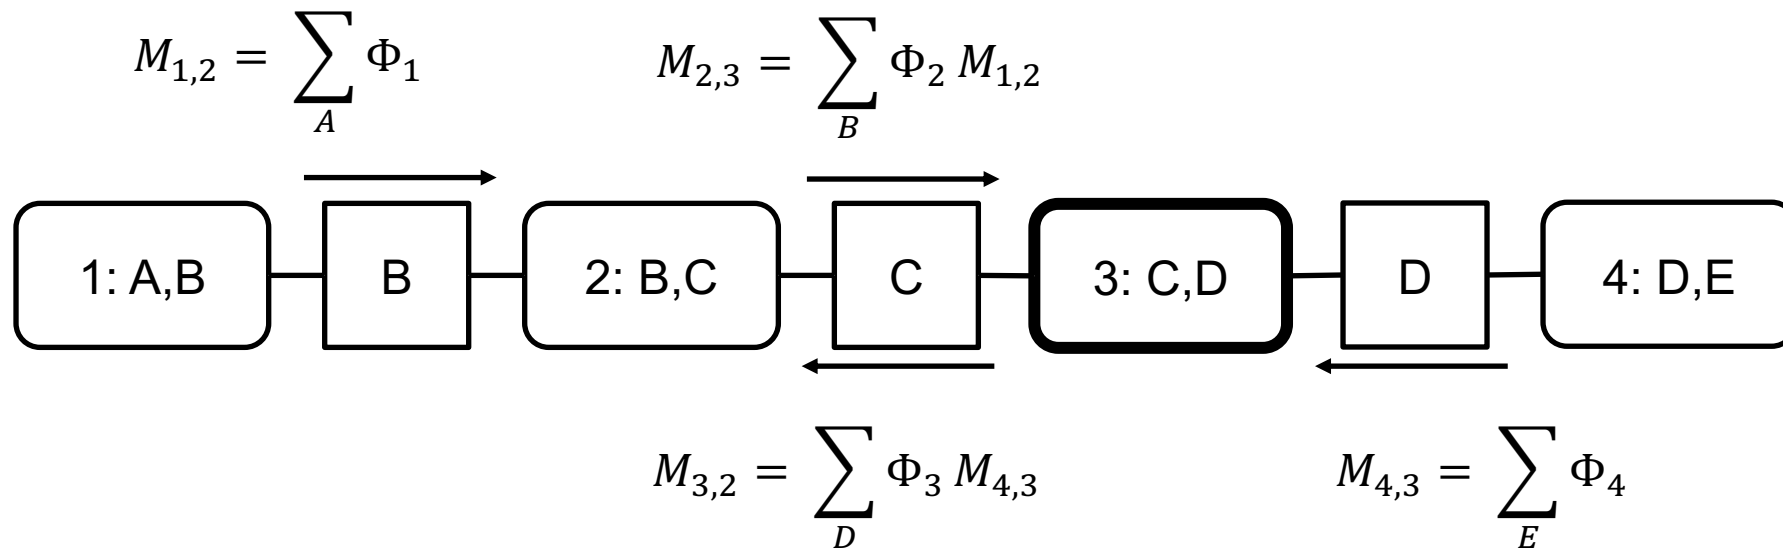

Message-passing: Example



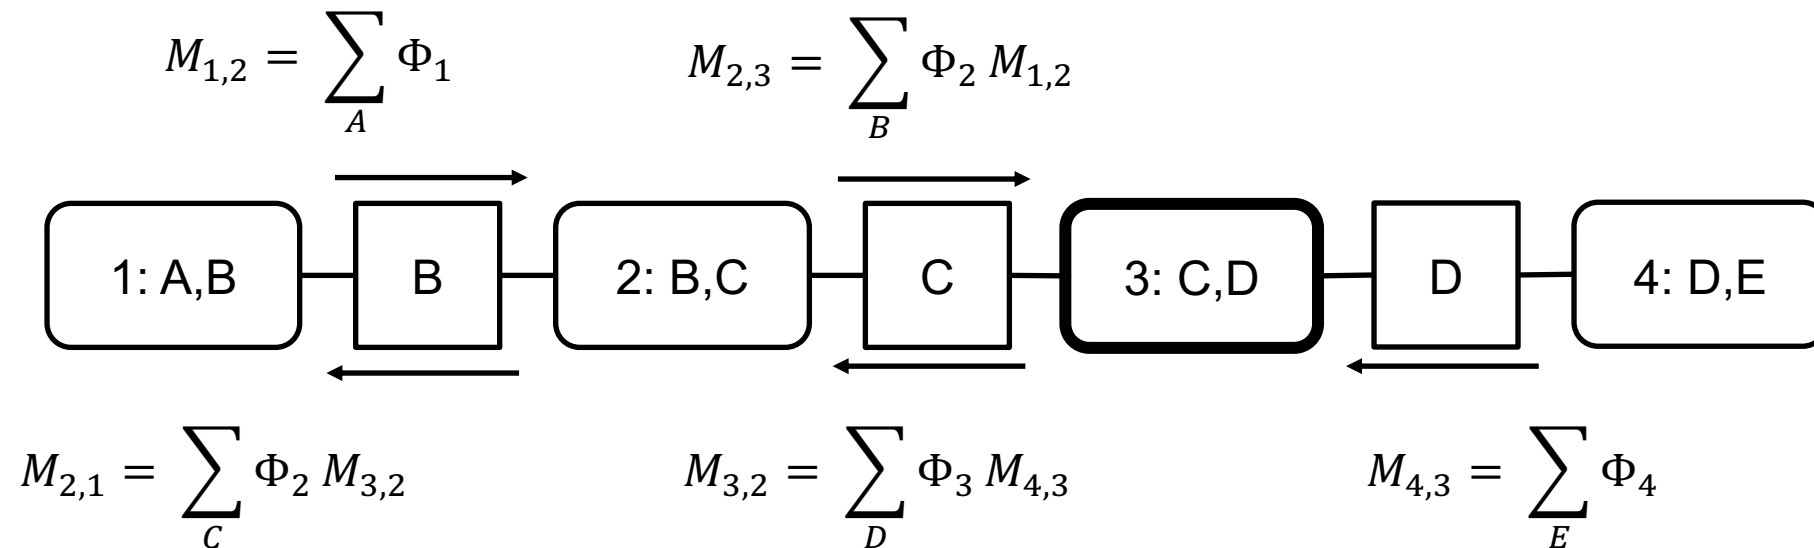

Message-passing: Example



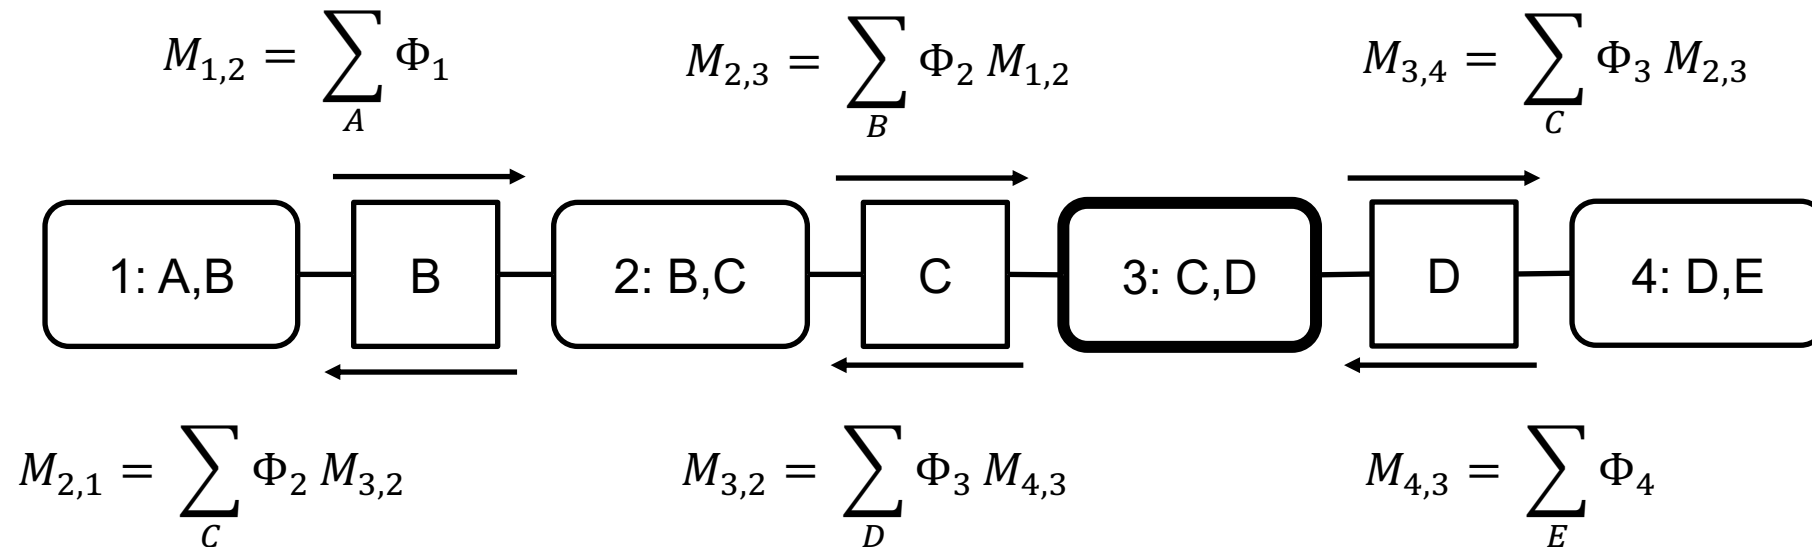

Message-passing: Example




Message-passing: Example



Message-passing: Example



Message-passing: Example



$$B_3(C, D) = \Phi_3 M_{2,3} M_{4,3}$$


$$= \Phi_3 \left(\sum_B \Phi_2 M_{1,2} \right) \sum_E \Phi_4$$

$$= \Phi_3 \left(\sum_B \Phi_2 \sum_A \Phi_1 \right) \sum_E \Phi_4$$



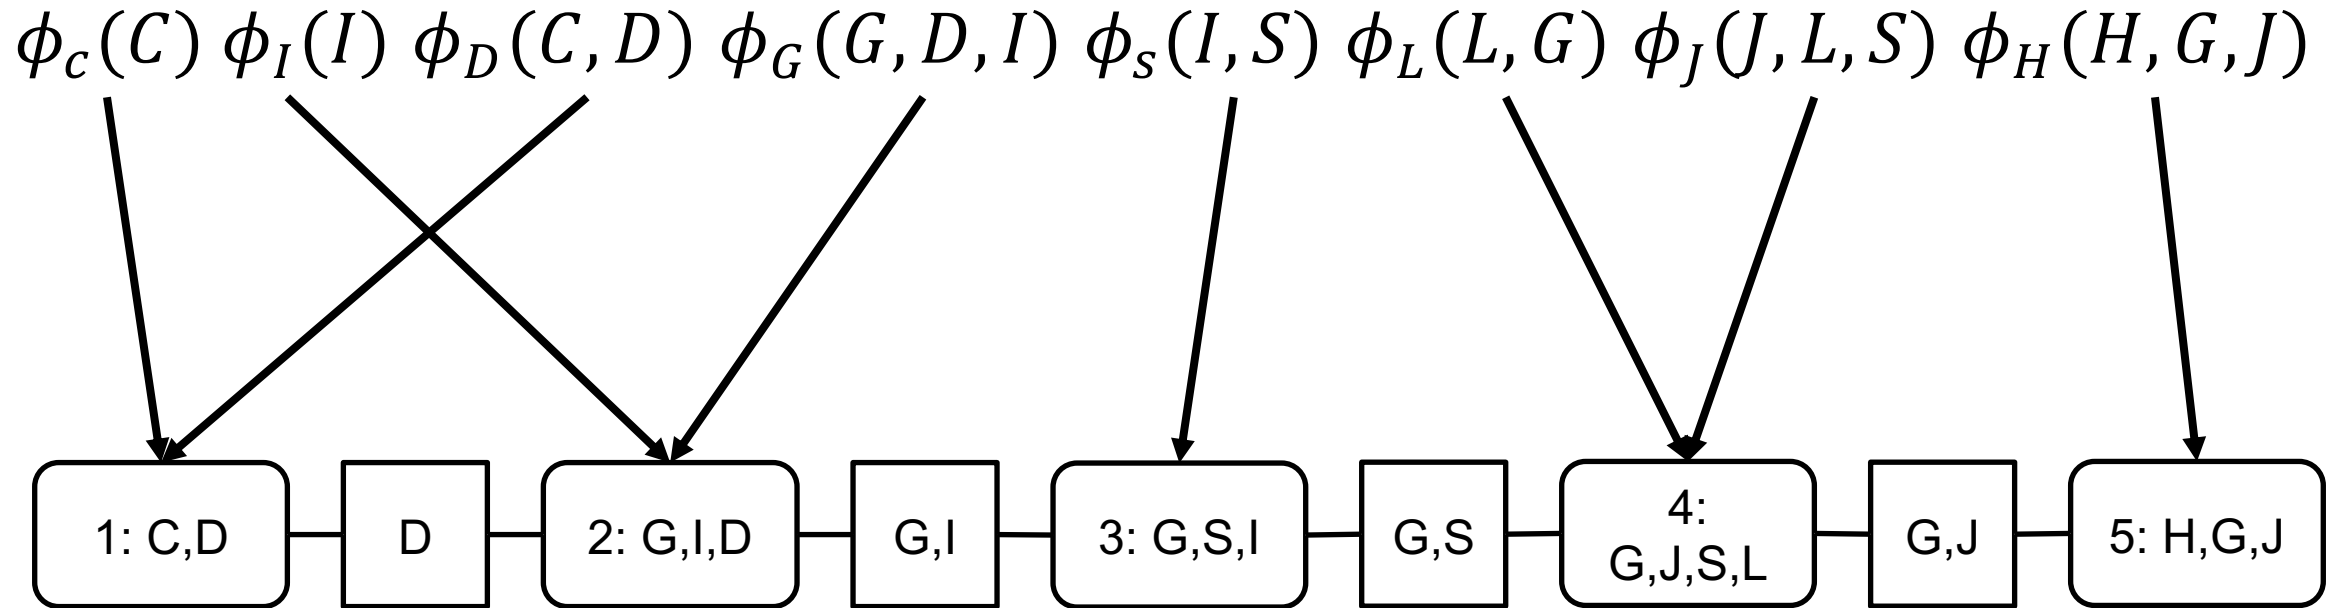
Product of all factors
marginalising hidden
variables in the correct order

Another Jointree Example

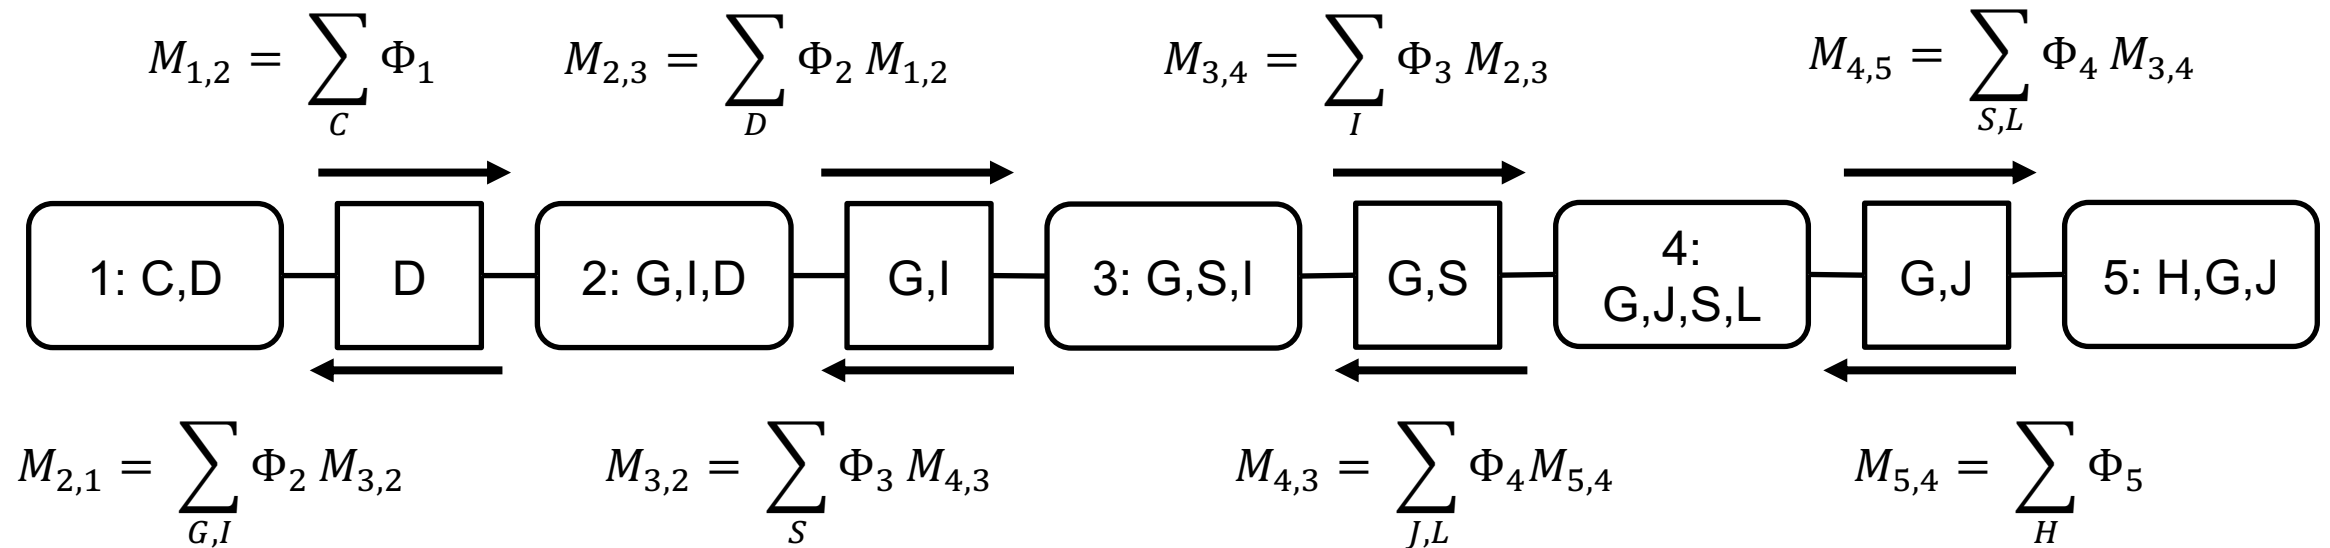
- Let's see another example with a larger Bayesian network
 - These are the factors
$$\phi_C(C), \phi_I(I), \phi_D(C, D),$$
$$\phi_G(G, D, I), \phi_S(I, S), \phi_L(L, G),$$
$$\phi_J(J, L, S), \phi_H(H, G, J)$$



Another Jointree Example



Another Jointree Example




Answering Queries

- For marginal probability queries on variables that appear together in a clique
 - Sum out irrelevant variables from any clique containing those variables
- For posterior marginal queries
 - We have evidence $E = e$ and query Q
- If Q and E appear in a same cluster
 - That is, $Q \cup E \subseteq C_i$ for some cluster C_i
 - We can eliminate entries that do not agree with evidence e
 - Sum out irrelevant variables and renormalize
- If Q does not appear in a cluster with E
 - Set evidence indicators in one or more clique containing E
 - Propagate messages along path to clique containing Q
 - Sum out irrelevant variables

Answering Queries

- Suppose we want to compute $P(C, a)$
 - We need set evidence and propagate some messages again
 - Let's now evidence indicators instead of elimination rows



$$M_{2,1} = \sum_C \Phi_2 M_{3,2}$$

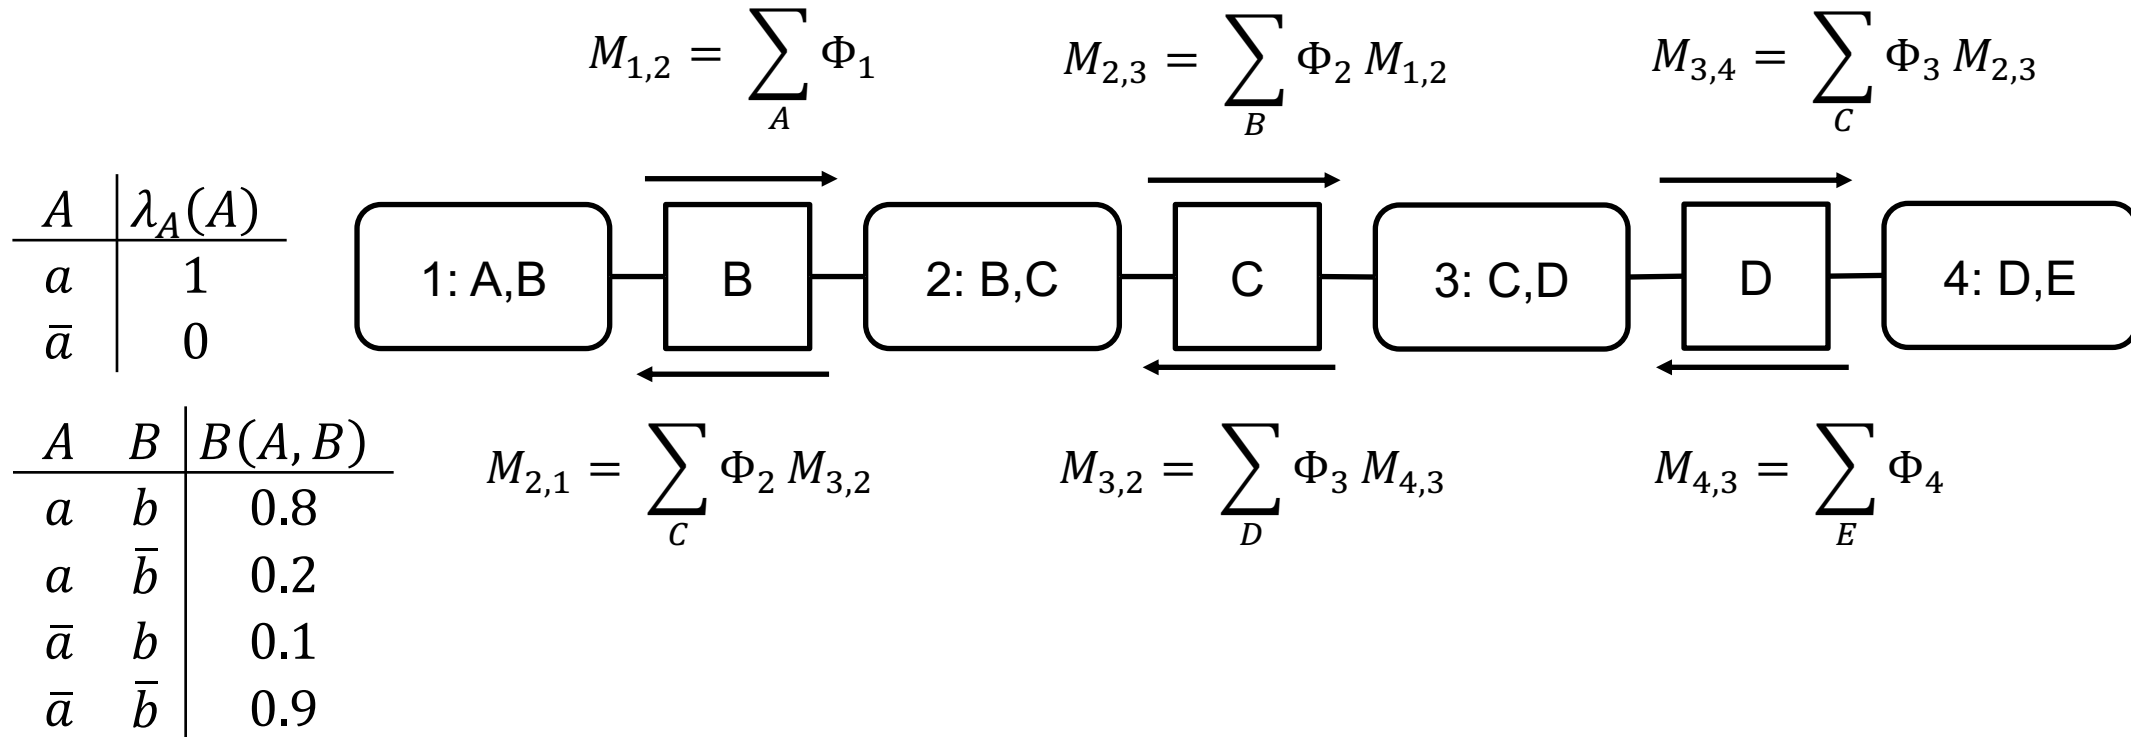
$$M_{3,2} = \sum_D \Phi_3 M_{4,3}$$

$$M_{4,3} = \sum_E \Phi_4$$

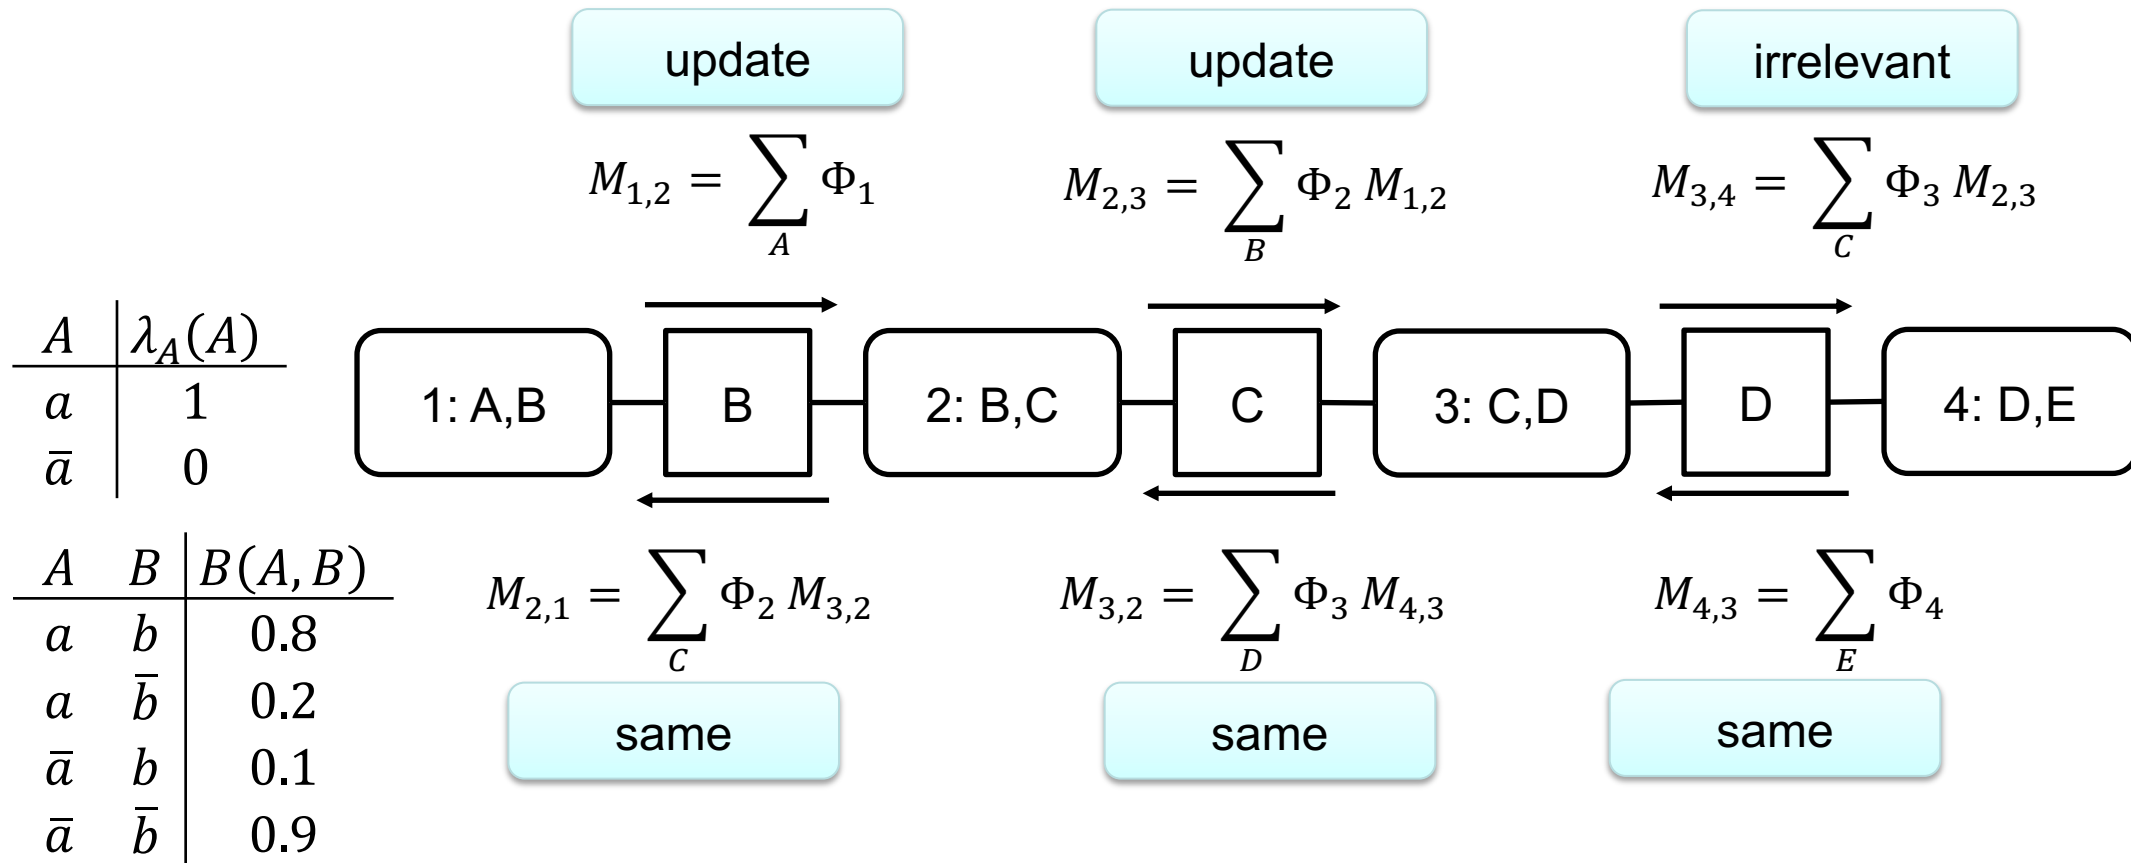
A	B	$B(A, B)$
a	b	0.8
a	\bar{b}	0.2
\bar{a}	b	0.1
\bar{a}	\bar{b}	0.9

Answering Queries

- Suppose we want to compute $P(C, a)$
 - We need set evidence and propagate some messages again
 - Let's now evidence indicators instead of elimination rows



Answering Queries



The Jointree Algorithm

- There are two main methods for propagating messages in a jointree, known as
 - The Shenoy-Shafer architecture and
 - The Hugin architecture
- The methods differ in both their space and time complexity
 - The Shenoy-Shafer architecture generally require less space but more time on an arbitrary jointree

The Shenoy-Shafer Architecture

- Evidence e is entered into the jointree through evidence indicators
- A cluster is selected as the root
- Message propagation in two phases
 - Inward: toward the root
 - Outward: away from the root
- Inward phase is also known as the *collect* or *pull* phase
 - The outward phase is known as the *distribute* or *push* phase
- Node i send a message to j only when it has received messages from all other neighbors k

- A message from node i to j is a factor

$$M_{ij} \stackrel{\text{def}}{=} \sum_{\mathcal{C}_i \setminus \mathcal{S}_{ij}} \Phi_i \prod_{k \neq j} M_{ki}$$

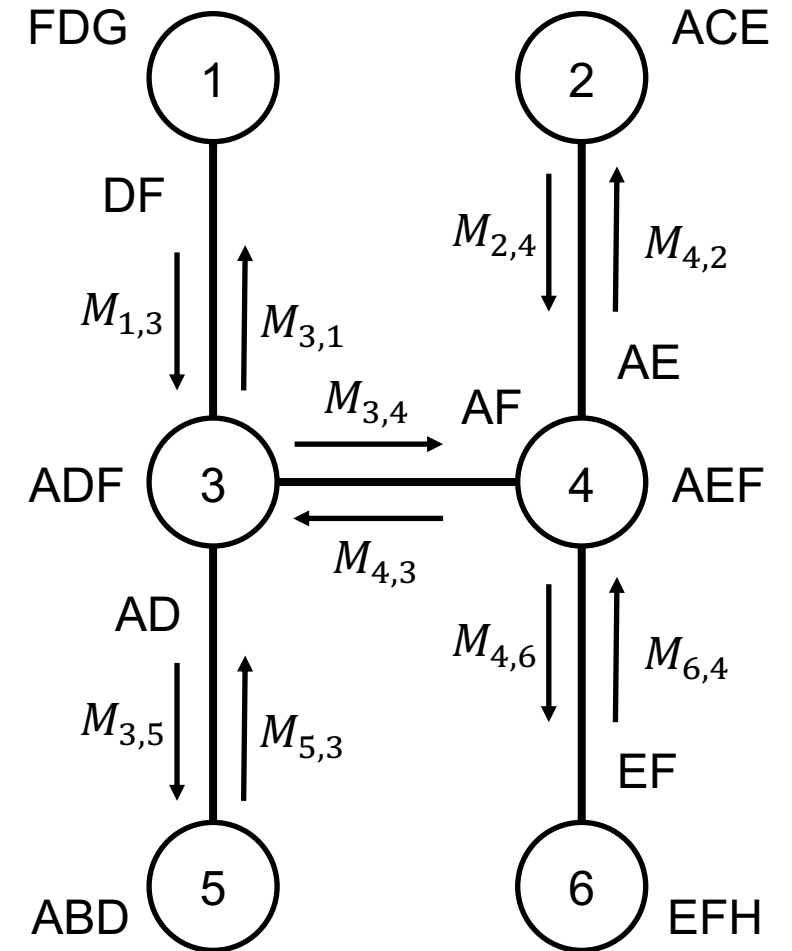
- Φ_i is the multiplication of all factors associated with node i (including evidence indicators)
- Once finished, we have the following for each cluster i in the jointree

$$P(\mathcal{C}_i) = \Phi_i \prod_k M_{ki}$$

- We can compute the joint marginal for any subset of variables that is included in a cluster

Shenoy-Shafer: Space

- We need two factors for each separator S_{ij}
 - One factor stores the message from cluster i to cluster j
 - The other stores the message from j to i
- There is no need to construct a factor over all cluster variables
 - The space complexity is not exponential in the size of jointree clusters
 - But only in the size of jointree separators

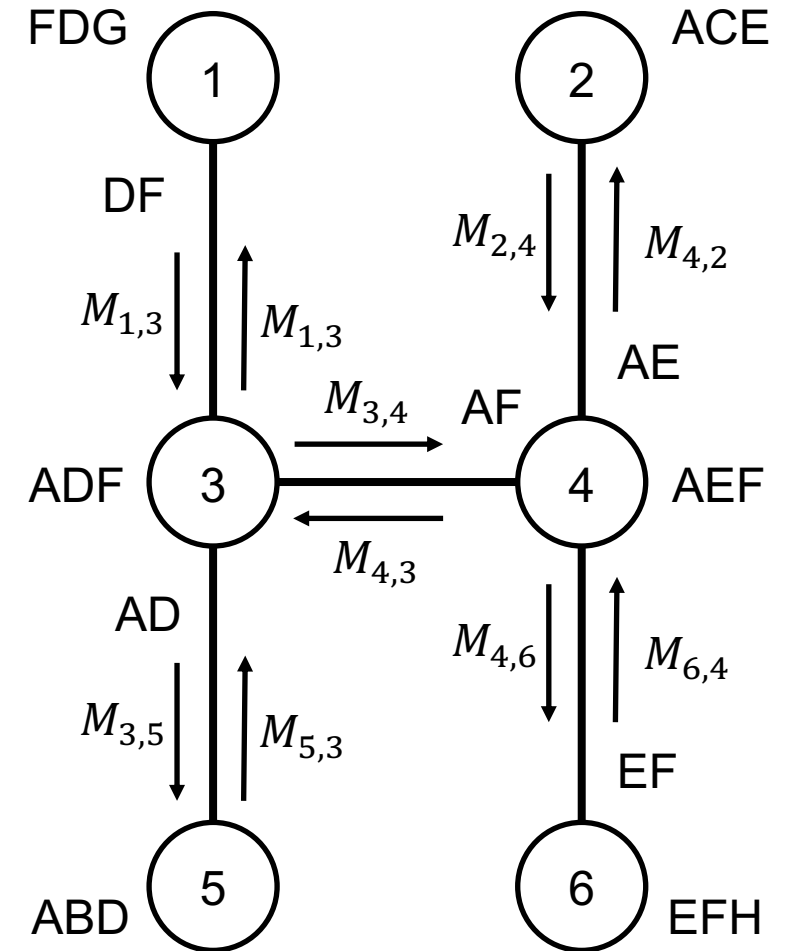


The Hugin Architecture

- Evidence e is entered into the jointree through evidence indicators
- A cluster is selected as the root
- Message propagation in two phases
 - Inward: toward the root
 - Outward: away from the root
- Inward phase is also known as the *collect* or *pull* phase
 - The outward phase is known as the *distribute* or *push* phase
- Node i send a message to j only when it has received messages from all other neighbors k
- Each separator S_{ij} has a single factor Ψ_{ij}
 - With each entry initialized to 1
- Each cluster C_i has a factor Ψ_i
 - Which is initialized to $\Phi_i \prod_j \Psi_{ij}$
 - Where Φ_i is the product of factors (including evidence indicators) assigned to node i
- When node i is ready to send a message to node j , it does the following:
 - Saves the factor Ψ_{ij} into Ψ_{ij}^{old}
 - Computes the new factor $\Psi_{ij} \leftarrow \sum_{C_i \setminus S_{ij}} \Psi_i$
 - Computes the message: $M_{ij} = \Psi_{ij} / \Psi_{ij}^{old}$
 - Multiplies M_{ij} into node j : $\Psi_j \leftarrow \Psi_j M_{ij}$

The Hugin Architecture

- After the inward and outward passes, we have the following for each node i
 - $P(\mathcal{C}_i, \mathbf{e}) = \Psi_i$
- The Hugin propagation scheme also guarantees the following for each edge $i - j$:
 - $P(\mathcal{S}_{ij}, \mathbf{e}) = \Psi_{ij}$
- The space requirements for the Hugin architecture
 - One factor for each cluster and one factor for each separator
 - The cluster factors are usually much larger than separator factors. More space than Shenoy-Shafer
 - However, Hugin is faster. We do not need to multiply all the factors to compute a new message



Conclusion

- Factor elimination is an alternate view of variable elimination
 - We decompose the graphs eliminating one factor at a time, instead of one variable
 - This view provides an efficient approach to answer queries over cluster variables
- The key idea is to use a message-passing formulation
 - Saving the intermediate computations that can be used later to answer queries
 - Message-passing also forms the basis of approximate algorithm known as belief propagation
- Tasks
 - Read Chapter 7 from the textbook (Darwiche)