

COMP9418: Advanced Topics in Statistical Machine Learning

Belief Propagation

Instructor: Gustavo Batista


University of New South Wales

Introduction

- In this lecture, we discuss a class of approximate inference algorithms based on belief propagation
 - Belief propagation was introduced as an exact algorithm for networks with polytree structure
 - Later, applied to networks with arbitrary structure and produced high-quality approximations in certain cases
- We introduce generalization of the algorithm with a full spectrum of approximations
 - Belief propagation approximation at one end
 - Exact results at the other


Belief Propagation

- Belief propagation is a messaging-passing algorithm
 - Originally developed for exact inference in polytrees networks
 - A polytree is a network with only one undirected path between any two nodes
- The exact algorithm is a variation of the jointree
 - It computes $P(X, \mathbf{e})$ for every variable in the polytree
 - We discuss the approximate algorithm later on



Belief Propagation


- Suppose we want to apply the jointree algorithm under evidence $E = \text{true}$
 - In this case, we can create a “special” jointree has the same structure as the polytree
 - A node i in the jointree has cluster $\mathcal{C}_i = XU$, where U are the parents of X
 - Edge $U \rightarrow X$ in the jointree has separator $\mathcal{S}_{ij} = U$
- Therefore
 - Jointree width equals polytree treewidth
 - Each jointree message is over a single variable



Belief Propagation

- *Belief propagation* is the jointtree algorithm under these circumstances
 - Messages are notated differently based on the polytree
 - Message from node U to child X is denoted $\pi_X(U)$ (*causal support*)
 - Messages from node Y to parent X is denoted $\lambda_Y(X)$ (*diagnostic support*)
- The joint marginal for the family of variable X with parents U_i and children Y_i is given by


$$P(X\mathbf{U}) = \phi_X(X, \mathbf{U}) \prod_i \pi_X(U_i) \prod_j \lambda_{Y_j}(X)$$



Belief Propagation with Evidence

- In the presence of evidence, the belief propagation uses an evidence indicator $\lambda_e(X)$
 - $\lambda_e(x) = 1$ if x is consistent with the evidence e and zero otherwise
 - We can rewrite the joint marginal for the family of variable X with parents U_i , children Y_i and evidence e as

$$P(X\mathbf{U}, \mathbf{e}) = \lambda_e(X) \phi_X(X, \mathbf{U}) \prod_i \pi_X(U_i) \prod_j \lambda_{Y_j}(X)$$



Belief Propagation with Evidence


- Using this notation, diagnostic messages can be defined as

$$\lambda_X(U_i) = \sum_{X \setminus U \setminus \{U_i\}} \lambda_e(X) \phi_X(X, \mathbf{U}) \prod_{k \neq i} \pi_X(U_k) \prod_j \lambda_{Y_j}(X)$$

- And causal messages as

$$\pi_{Y_j}(X) = \sum_U \lambda_e(X) \phi_X(X, \mathbf{U}) \prod_i \pi_X(U_i) \prod_{k \neq j} \lambda_{Y_k}(X)$$

- A node can send a message to a neighbour only after it has received messages from all other neighbours



Belief Propagation with Evidence


- Using this notation, diagnostic messages can be defined as

$$\lambda_X(U_i) = \sum_{X \setminus U_i} \lambda_e(X) \phi_X(X, \mathbf{U}) \prod_{k \neq i} \pi_X(U_k) \prod_j \lambda_{Y_j}(X)$$

- And causal messages as

$$\pi_{Y_j}(X) = \sum_U \lambda_e(X) \phi_X(X, \mathbf{U}) \prod_i \pi_X(U_i) \prod_{k \neq j} \lambda_{Y_k}(X)$$

- A node can send a message to a neighbour only after it has received messages from all other neighbours



Belief Propagation with Evidence

- When a node has a single neighbour, it can immediately send a message to that neighbour

- This includes a leaf node X with a single parent U


$$\lambda_X(U) = \sum_X \lambda_e(X) \phi_X(X, U)$$

- And a root node X with a single child Y

$$\pi_Y(X) = \lambda_e(X) \phi_X(X, U)$$

- These are the base cases for belief propagation

- These messages can be computed immediately as they do not depend on any other messages
- Typically, messages are first propagated toward a root and then pushed away from root



Belief Propagation: Example

$$P(B, C, D, \mathbf{e}) = \phi_D(D, B, C) \pi_D(B) \pi_D(C) \lambda_E(D) \lambda_F(D)$$

$\mathbf{e}: \{E = \text{true}\}$

A	$\phi_A(A)$
a	.01
\bar{a}	.99

C	$\phi_C(C)$
c	.001
\bar{c}	.999

B	C	D	$\phi_D(D, B, C)$
b	c	d	.99
b	c	\bar{d}	.01
b	\bar{c}	d	.90
b	\bar{c}	\bar{d}	.10
\bar{b}	c	d	.95
\bar{b}	c	\bar{d}	.05
\bar{b}	\bar{c}	d	.01
\bar{b}	\bar{c}	\bar{d}	.99

A	$\pi_B(A)$
a	.01
\bar{a}	.99

C	$\pi_D(C)$
c	.001
\bar{c}	.999



B	C	D	$P(B, C, D, \mathbf{e})$
b	c	d	1.7731×10^{-6}
b	c	\bar{d}	5.9700×10^{-9}
b	\bar{c}	d	1.6103×10^{-3}
b	\bar{c}	\bar{d}	5.9640×10^{-5}
\bar{b}	c	d	8.5330×10^{-4}
\bar{b}	c	\bar{d}	1.4970×10^{-5}
\bar{b}	\bar{c}	d	8.9731×10^{-3}
\bar{b}	\bar{c}	\bar{d}	2.9611×10^{-1}

B	$\pi_D(B)$
b	.00199
\bar{b}	.99801

D	F	$\phi_F(F, D)$
d	f	.2
d	\bar{f}	.8
\bar{d}	f	.1
\bar{d}	\bar{f}	.9

D	$\lambda_E(D)$
d	.9
\bar{d}	.3

D	$\lambda_F(D)$
d	1
\bar{d}	1



Belief Propagation: Example


- We can use $P(B, C, D, \mathbf{e})$ to compute marginals for the variables B , C and D . For instance
 - We can also compute the joint marginal for C once we compute the message from D to C
 - To compute conditional marginals, we simply normalize joint marginals
- Another approach is to use a constant η that normalizes the factor to sum to one

C	$P(C, \mathbf{e})$
c	.0009
\bar{c}	.3067

$$P(X\mathbf{U}|\mathbf{e}) = \eta \lambda_{\mathbf{e}}(X) \phi_X(X, \mathbf{U}) \prod_i \pi_X(U_i) \prod_j \lambda_{Y_j}(X)$$


$$\lambda_X(U_i) = \eta \sum_{X\mathbf{U} \setminus \{U_i\}} \lambda_{\mathbf{e}}(X) \phi_X(X, \mathbf{U}) \prod_{k \neq i} \pi_X(U_k) \prod_j \lambda_{Y_j}(X)$$

$$\pi_{Y_j}(X) = \eta \sum_{\mathbf{U}} \lambda_{\mathbf{e}}(X) \phi_X(X, \mathbf{U}) \prod_i \pi_X(U_i) \prod_{k \neq j} \lambda_{Y_k}(X)$$




Belief Propagation in Connected Networks

- Belief propagation was designed as an exact algorithm for polytrees
 - However, it was later applied to connected networks
- This application poses some difficulties
 - A message can be sent from X to Y only when X has received all messages from other neighbours
 - The correctness of belief propagation depends on the underlying polytree
- The results can be incorrect if applied to connected networks
 - The algorithm is no longer always correct
 - But can still provide some high-quality approximations in many cases
- In the figure, after node E send a message to C no other message can be propagated
 - Since each is dependent on others that are waiting to be propagated




Iterative Belief Propagation (IBP)

- *Iterative or Loopy Belief Propagation* assumes some initial value to each message in the network
 - Given these initial values, each node is ready to send a message to each of its neighbours
 - At each iteration t , every node X send a message to its neighbours using the messages received from its other neighbours in $t - 1$
- The algorithm iterates until message convergence
 - The value of messages at the current iteration are within some threshold from their values at the previous iteration
 - When IBP converges, the values of the messages at convergence are called *fixed point*
 - IBP may have multiple fixed points on a given network



Message Schedule

- For some networks, IBP can oscillate and never converge
- The convergence rate can depend on the order the messages are propagated, which is known as *message schedule*
 - Parallel schedule: the order of the messages does not affect the algorithm
 - Sequential schedule: messages are propagated as soon as they are computed
- Sequential schedules are flexible in when and how quickly information is propagated
- Although one schedule may converge and other may not, all schedules have the same fixed points



Parallel Iterative Belief Propagation

$t \leftarrow 0$

initialize all messages

while messages have not converged **do**

$t \leftarrow t + 1$

for each node X with parents \mathbf{U} **do**

for each parent U_i **do**

$$\lambda_X^t(U_i) = \eta \sum_{X \setminus \{U_i\}} \lambda_e(X) \phi_X(X, \mathbf{U}) \prod_{k \neq i} \pi_X^{t-1}(U_k) \prod_j \lambda_{Y_j}^{t-1}(X)$$

for each child Y_j **do**

$$\pi_{Y_j}^t(X) = \eta \sum_{\mathbf{U}} \lambda_e(X) \phi_X(X, \mathbf{U}) \prod_i \pi_X^{t-1}(U_i) \prod_{k \neq j} \lambda_{Y_k}^{t-1}(X)$$

return $\beta(X\mathbf{U}) = P(X\mathbf{U}|\mathbf{e}) = \eta \lambda_e(X) \phi_X(X, \mathbf{U}) \prod_i \pi_X^t(U_i) \prod_j \lambda_{Y_j}^t(X)$

The Kullback-Leibler Divergence

- The *Kullback-Leibler divergence*, known as *KL divergence*, between two distributions P and P' conditioned on \mathbf{e}

$$KL(P'(X|\mathbf{e}), P(X|\mathbf{e})) \stackrel{\text{def}}{=} \sum_x P'(\mathbf{x}|\mathbf{e}) \log \frac{P'(\mathbf{x}|\mathbf{e})}{P(\mathbf{x}|\mathbf{e})}$$

- $KL(P'(X|\mathbf{e}), P(X|\mathbf{e}))$ is non-negative and equal to zero if and only if $P'(X|\mathbf{e})$ and $P(X|\mathbf{e})$ are equivalent
 - However, KL divergence is not a true distance since it is not symmetric. In general
$$KL(P'(X|\mathbf{e}), P(X|\mathbf{e})) \neq KL(P(X|\mathbf{e}), P'(X|\mathbf{e}))$$
 - We say we are weighting the KL divergence by the approximate distribution
 - This variation has some useful computational properties

Optimizing KL Divergence

- The approximate inference can be posed as an optimization problem
 - The goal is to search for an approximate distribution P' that minimizes KL divergence with P
 - We can assume a parametrized form for P' and search for the best instance, i.e., the best set of parameters
- The Iterative Belief Propagation algorithm presented before assumes that the approximate distribution $P'(X)$ factors as

$$P'(X|e) = \prod_{XU} \frac{P'(XU|e)}{\prod_{U \in U} P'(U|e)}$$


- XU ranges over the families of the network N
- U ranges over nodes that appear as parents in N

Optimizing KL Divergence

- The approximate distribution $P'(X)$ factors as

$$P'(X|e) = \prod_{XU} \frac{P'(XU|e)}{\prod_{U \in U} P'(U|e)}$$

- XU ranges over the families of the network N
 - U ranges over nodes that appear as parents in N
- Some observations about this assumption
 - This choice of $P'(X|e)$ is expressive enough to describe distributions induced by polytree networks
 - If the network N is a polytree then $P(X|e)$ does factor according to this equation (see figure for an example)
 - If N is not a polytree, then we are trying to fit $P(X|e)$ into an approximation $P'(X|e)$ as if it were generated by a polytree



$$\frac{P(A, B, C, D, E, F) = P(A)P(C)P(B, A) P(D, B, C)P(E, D)P(F, D)}{P(A)P(B)P(C)P(D)P(D)}$$

Optimizing KL Divergence

- The previous correspondence that IBP fixed points are stationary points of the KL divergence
 - They may or may not be local minima
 - When IBP performs well, it often has fixed points that are minima of the KL divergence
 - Otherwise, we need to seek approximations P' whose factorizations are more expressive than the polytree-based factorization
- If we do not insist on marginals being over families and individual variables, we can have a more general form that covers every distribution

Generalized Belief Propagation

- We saw in the previous lecture that a network can be factorized according to this expression if
 - \mathcal{C} corresponds to the clusters of a jointree
 - \mathcal{S} corresponds to the separators
- If we base our factorization in a jointree
 - Solving the previous optimization problem yields the same update equations of the jointree algorithm
- Therefore, the factorizations used by IBP and the factorization based on jointrees can be viewed as two extremes
 - One efficient but approximate
 - The other expensive but exact

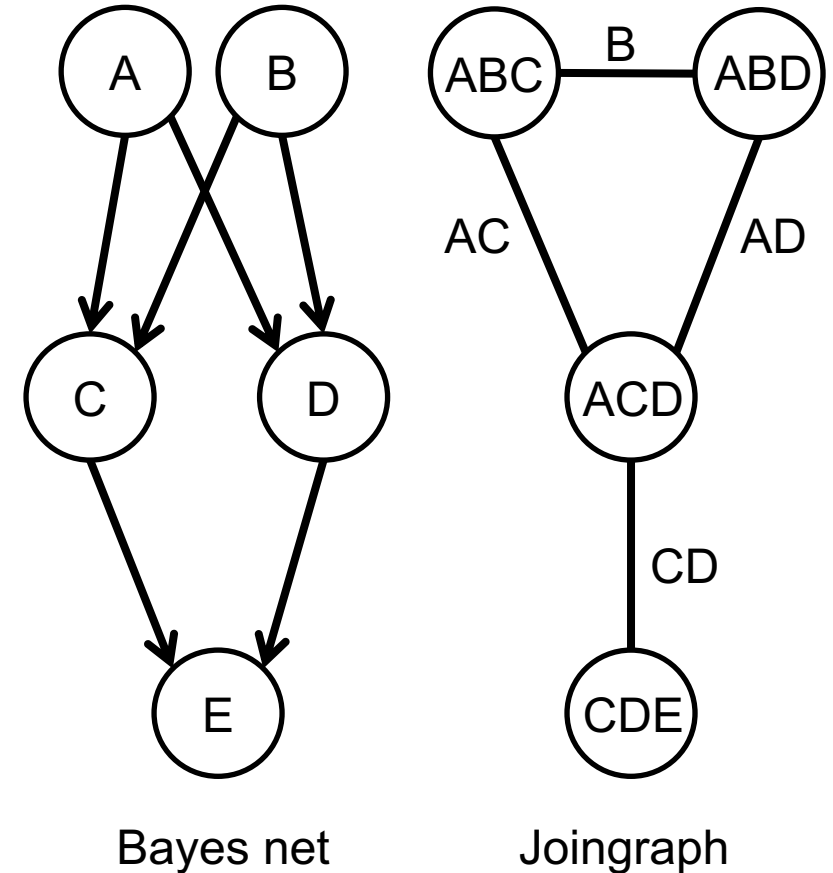
$$P'(\mathbf{X}|\mathbf{e}) = \frac{\prod_{\mathcal{C}} P'(\mathcal{C}|\mathbf{e})}{\prod_{\mathcal{S}} P'(\mathcal{S}|\mathbf{e})}$$

Joingraphs

- There is a spectrum of factorizations that fall in between these two extremes
 - This allows a trade-off quality and efficiency
 - The notion of joingraph is one way to obtain such a spectrum
- *Joingraphs* are generalizations of jointrees
 - They can be used to obtain factorizations according to $P'(\mathbf{X}|\mathbf{e}) = \frac{\prod_{\mathcal{C}} P'(\mathcal{C}|\mathbf{e})}{\prod_{\mathcal{S}} P'(\mathcal{S}|\mathbf{e})}$
 - They are used to formulate a message-passing algorithm similar to IBF, known as *iterative joingraph propagation*


Joingraphs

- A *joingraph* G for a network N is a graph where nodes i are labelled by cluster \mathcal{C}_i , and edges $i - j$ are labelled by separators \mathcal{S}_{ij} . Moreover, G satisfies the following properties
 - Clusters \mathcal{C}_i and separators \mathcal{S}_{ij} are sets of nodes from N
 - Each factor in N must appear in some cluster \mathcal{C}_i
 - If a variable X appears in two clusters \mathcal{C}_i and \mathcal{C}_j , then there exists a unique path connecting i and j in the joingraph such that X appears in every cluster and separator on that path
 - For every edge $i - j$ in the joingraph, $\mathcal{S}_{ij} \subseteq \mathcal{C}_i \cap \mathcal{C}_j$



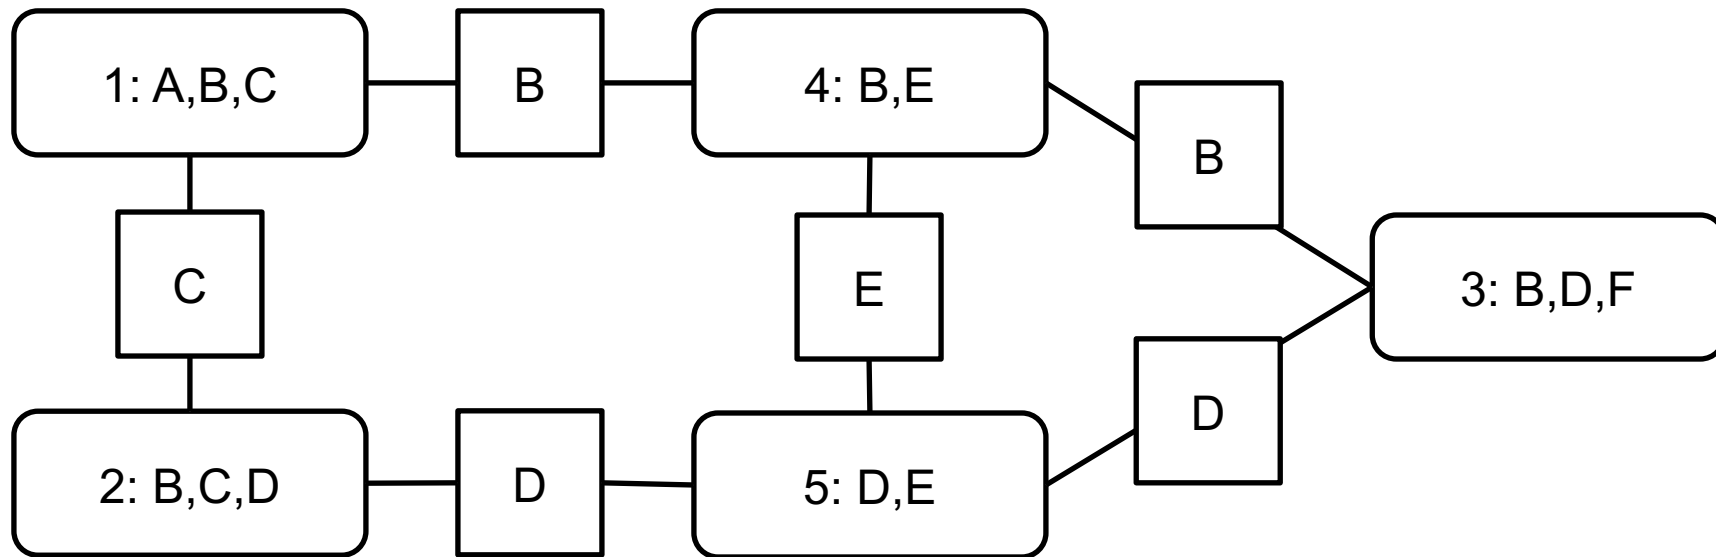
Jointrees and Joingraphs

- We can think of a jointgraph as a way of relaxing some constraints of jointtrees
 - In a jointtree, if two clusters C_i and C_j share a set of variables X then every cluster and separator on the path connecting C_i and C_j must contain X
 - In a joingraph, we assert each variable $X \in X$ be contained in clusters and separators of some path connecting C_i and C_j
 - We do not require separators S_{ij} to be precisely the intersection of C_i and C_j , as in the case of jointtrees




Valid Joingraph?

$\phi_1(A, B, C), \phi_2(B, C), \phi_3(B, D), \phi_4(D, E), \phi_5(B, E), \phi_6(B, D), \phi_7(B, D, F)$




Valid Joingraph?

$\phi_1(A, B, C), \phi_2(B, C), \phi_3(B, D), \phi_4(D, E), \phi_5(B, E), \phi_6(B, D), \phi_7(B, D, F)$



Valid Joingraph?

$\phi_1(A, B, C), \phi_2(B, C), \phi_3(B, D), \phi_4(D, E), \phi_5(B, E), \phi_6(B, D), \phi_7(B, D, F)$




Joingraph Factorization

- A joingraph induces an approximate factorization

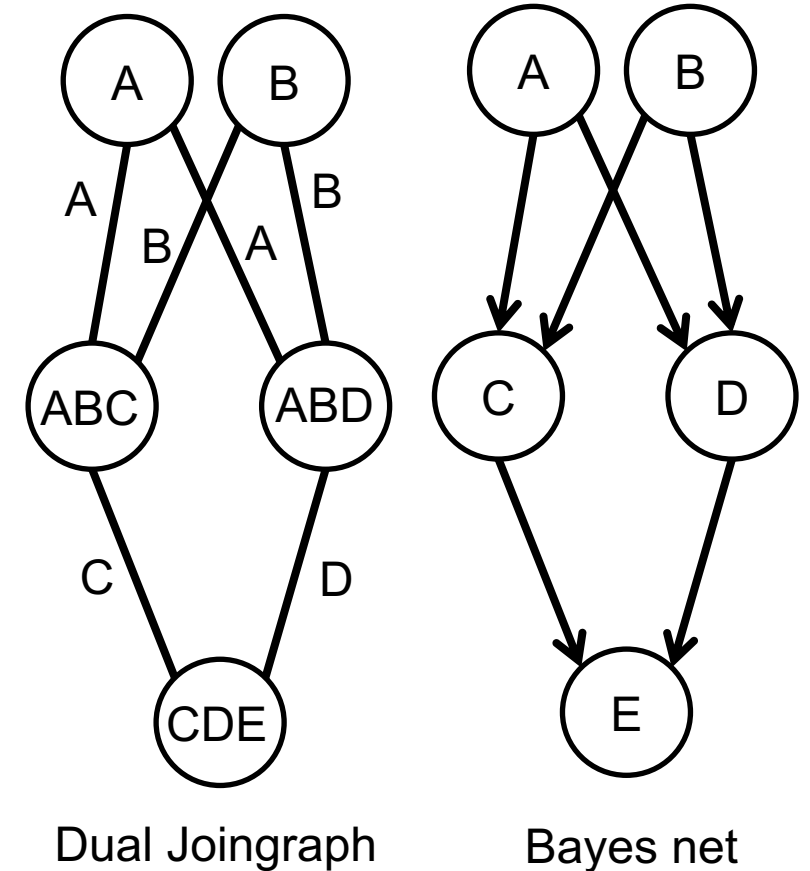
$$P'(\mathbf{X}|\mathbf{e}) = \frac{\prod_i P'(\mathbf{C}_i|\mathbf{e})}{\prod_{ij} P'(\mathbf{S}_{ij}|\mathbf{e})}$$

- When the joingraph corresponds to a jointree, the factorization is exact



Dual Joingraph


- A *dual joingraph* is a special joingraph whose factorization reduces to the one used by IBP
- A dual joingraph G for a network N is obtained as follows
 - G has the same undirected structure as N
 - For each family XU in N , the corresponding node i in G has the cluster $\mathbf{C}_i = XU$
 - For each $U \rightarrow X$ in N , the corresponding edge $i - j$ in G has separator $\mathbf{S}_{ij} = U$




Factorization Comparison

- Dual jointgraph (approximate, same as IBP)

$$P'(X|e) = \frac{P'(A|e)P'(B|e)P'(A, B, C|e)P'(A, B, D|e)P'(C, D, E|e)}{P'(A|e)^2 P'(B|e)^2 P'(C|e)P'(D|e)}$$



Dual Joingraph



Bayes net

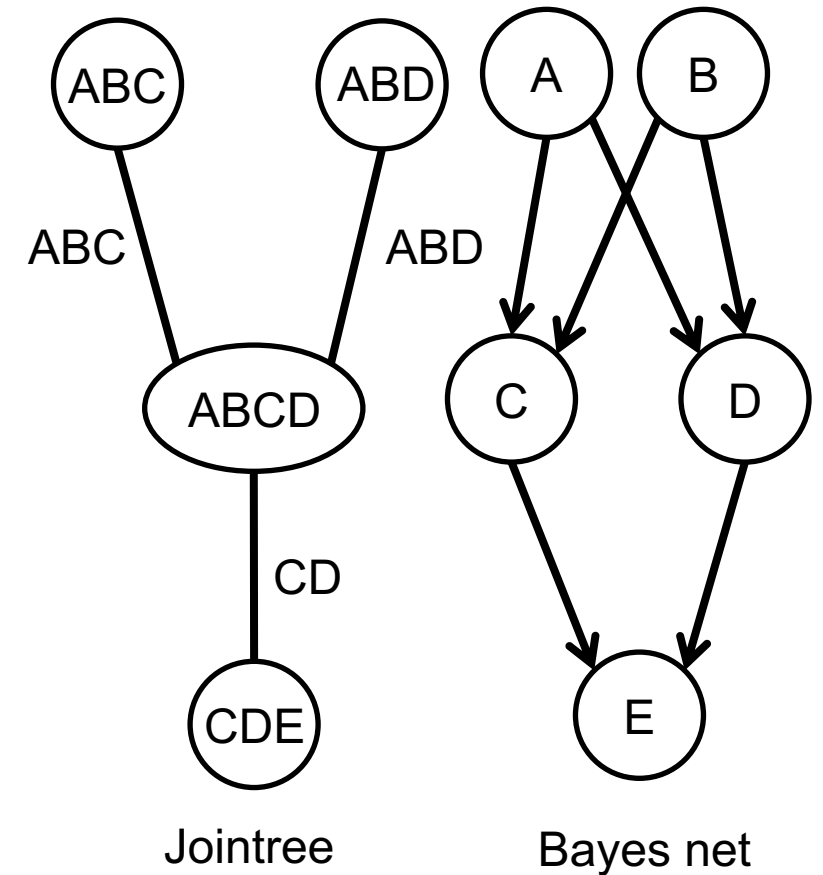
Factorization Comparison

- Dual jointgraph (approximate, same as IBP)

$$P'(X|e) = \frac{P'(A|e)P'(B|e)P'(A, B, C|e)P'(A, B, D|e)P'(C, D, E|e)}{P'(A|e)^2 P'(B|e)^2 P'(C|e)P'(D|e)}$$

- Jointree (exact)

$$P'(X|e) = \frac{P'(A, B, C|e)P'(A, B, D|e)P'(A, B, C, D|e)P'(C, D, E|e)}{P'(A, B, C|e)P'(A, B, D|e)P'(C, D|e)}$$



Factorization Comparison

- Dual jointgraph (approximate, same as IBP)


$$P'(X|e) = \frac{P'(A|e)P'(B|e)P'(A, B, C|e)P'(A, B, D|e)P'(C, D, E|e)}{P'(A|e)^2 P'(B|e)^2 P'(C|e)P'(D|e)}$$

- Jointree (exact)

$$P'(X|e) = \frac{P'(A, B, C|e)P'(A, B, D|e)P'(A, B, C, D|e)P'(C, D, E|e)}{P'(A, B, C|e)P'(A, B, D|e)P'(C, D|e)}$$

- Joingraph (trade complexity and quality)

$$P'(X|e) = \frac{P'(A, B, C|e)P'(A, B, D|e)P'(A, C, D|e)P'(C, D, E|e)}{P'(B|e)P'(A, C|e)P'(A, D|e)P'(C, D|e)}$$



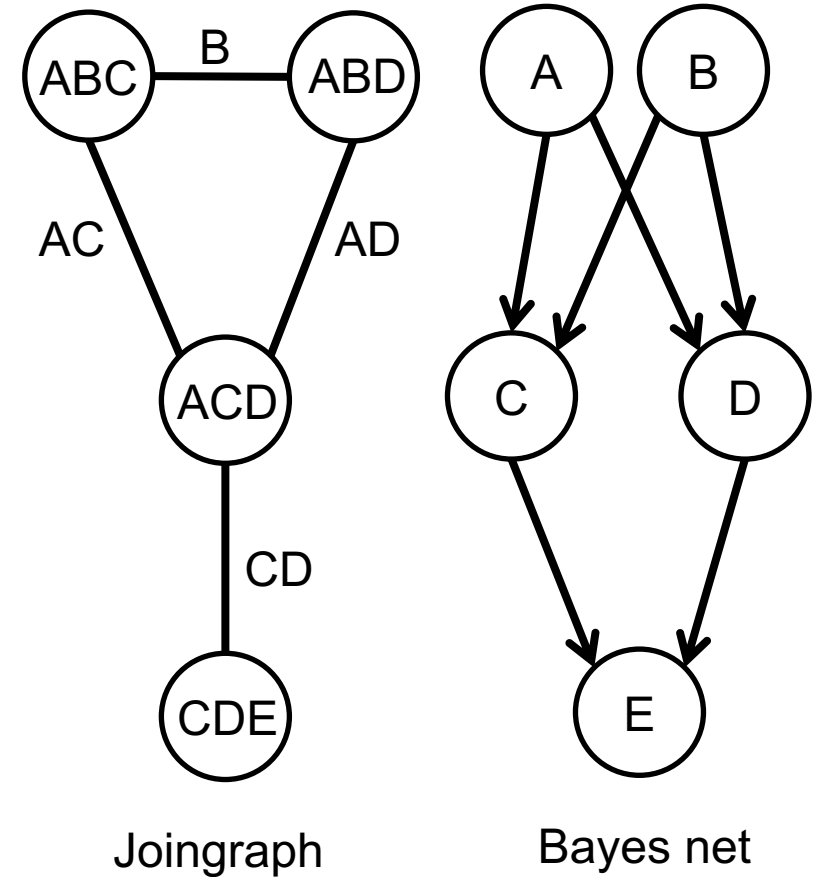
Joingraph

Bayes net

Factorization Comparison

- Joining graph (trade complexity and quality)

$$P'(X|e) = \frac{P'(A, B, C|e)P'(A, B, D|e)P'(A, C, D|e)P'(C, D, E|e)}{P'(B|e)P'(A, C|e)P'(A, D|e)P'(C, D|e)}$$



Iterative Joingraph Propagation

- Suppose we have a network N that induces a distribution P
 - And a corresponding joingraph that induces a factorization P'
 - Also, we want to compute cluster marginals $P'(\mathcal{C}_i|\mathbf{e})$ and separator marginals $P'(\mathcal{S}_{ij}|\mathbf{e})$ that minimize the KL divergence between $P(\mathbf{X}|\mathbf{e})$ and $P'(\mathbf{X}|\mathbf{e})$
- This optimization problem can be solved with a generalization of IBP called *interactive joingraph propagation* (IJGP)

Iterative Joingraph Propagation

- The algorithm starts assigning each network factor ϕ and evidence indicator λ_e to some cluster \mathcal{C}_i that contains variables in ϕ
 - All factors are associated to some cluster (no information loss)
 - No factor is present in more than one cluster (no overcounting of information)
- It propagates messages with the equations
 - $M_{ij} = \eta \sum_{\mathcal{C}_i \setminus \mathcal{S}_{ij}} \psi_i \prod_{k \neq j} M_{ki}$
 - where ψ_i is the product of all CPTs and evidence indicators assigned to cluster \mathcal{C}_i
 - M_{ij} is the message sent from cluster i to cluster j

Parallel Iterative Joingraph Propagation

$t \leftarrow 0$

initialize all messages

while messages have not converged **do**

$t \leftarrow t + 1$

for each joingraph edge $i - j$ **do**

$$M_{ij}^t = \eta \sum_{\mathbf{C}_i \setminus \mathbf{S}_{ij}} \psi_i \prod_{k \neq j} M_{ki}^{t-1}$$


$$M_{ji}^t = \eta \sum_{\mathbf{C}_i \setminus \mathbf{S}_{ij}} \psi_i \prod_{k \neq i} M_{kj}^{t-1}$$

return $\beta(\mathbf{C}_i) = \eta \psi_i \prod_k M_{ki}^t$ for each node i

Joingraph Example with Markov Nets

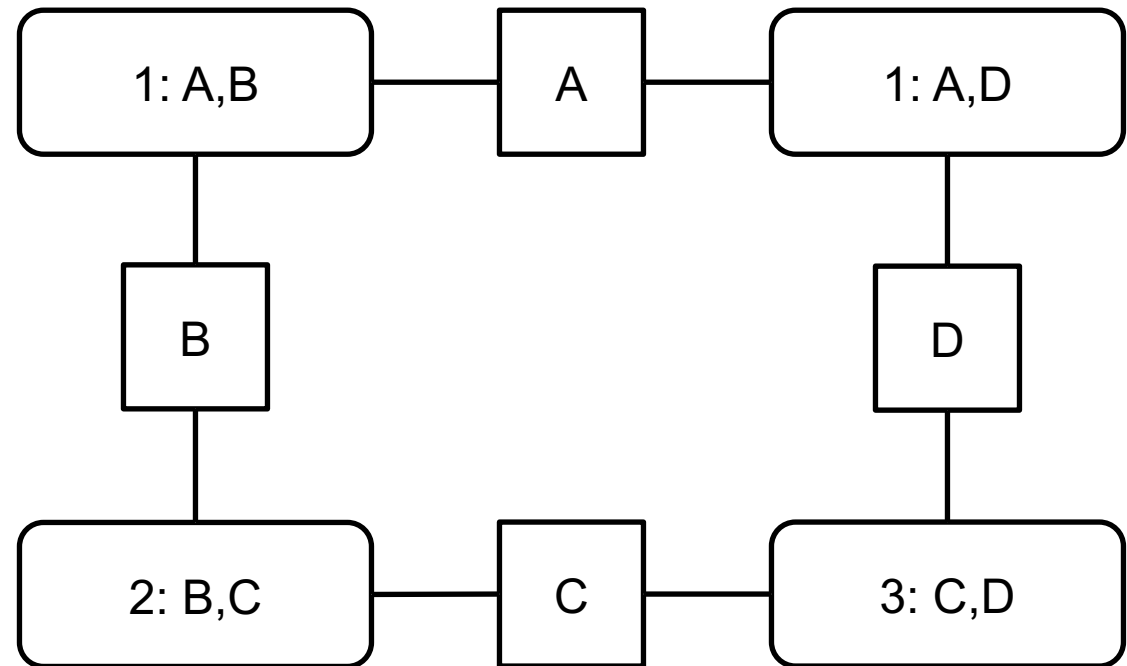
D	A	$\phi_4(D, A)$
d	a	100
d	\bar{a}	1
\bar{d}	a	1
\bar{d}	\bar{a}	100

A	B	$\phi_1(A, B)$
a	b	30
a	\bar{b}	5
\bar{a}	b	1
\bar{a}	\bar{b}	10



C	D	$\phi_3(C, D)$
c	d	1
c	\bar{d}	100
\bar{c}	d	100
\bar{c}	\bar{d}	1


B	C	$\phi_2(B, C)$
b	c	100
b	\bar{c}	1
\bar{b}	c	1
\bar{b}	\bar{c}	100



Joingraph Example with Markov Nets

D	A	$\phi_4(D, A)$
d	a	100
d	\bar{a}	1
\bar{d}	a	1
\bar{d}	\bar{a}	100

A	B	$\phi_1(A, B)$
a	b	30
a	\bar{b}	5
\bar{a}	b	1
\bar{a}	\bar{b}	10

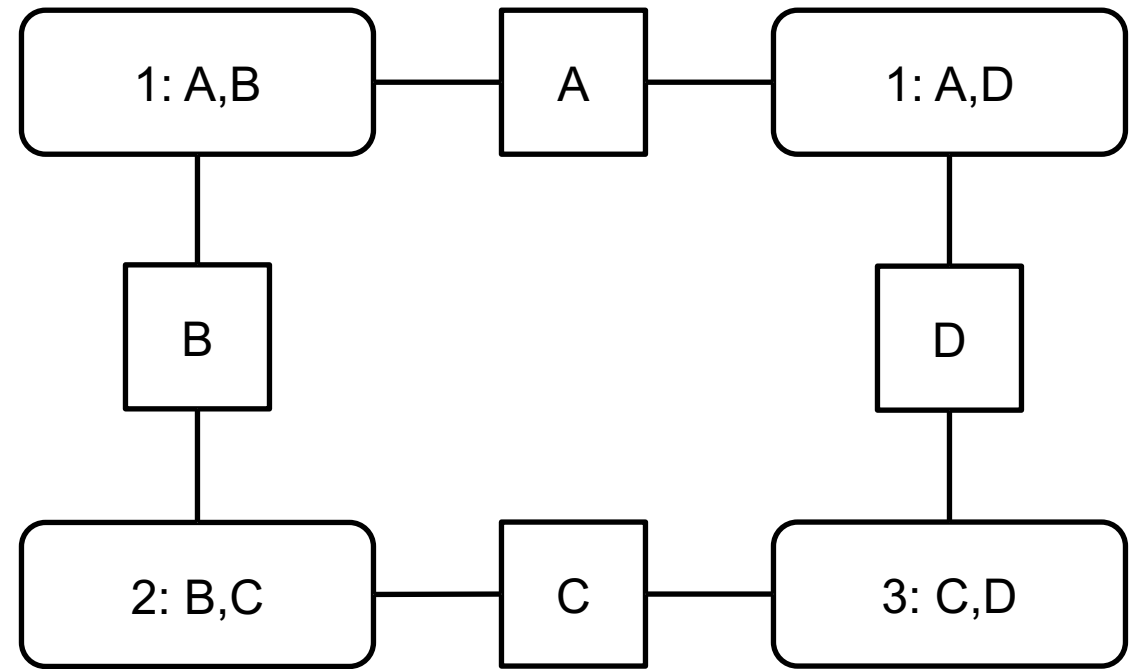


C	D	$\phi_3(C, D)$
c	d	1
c	\bar{d}	100
\bar{c}	d	100
\bar{c}	\bar{d}	1

B	C	$\phi_2(B, C)$
b	c	100
b	\bar{c}	1
\bar{b}	c	1
\bar{b}	\bar{c}	100

$$\psi_1(A, B) = \phi_1(A, B)$$

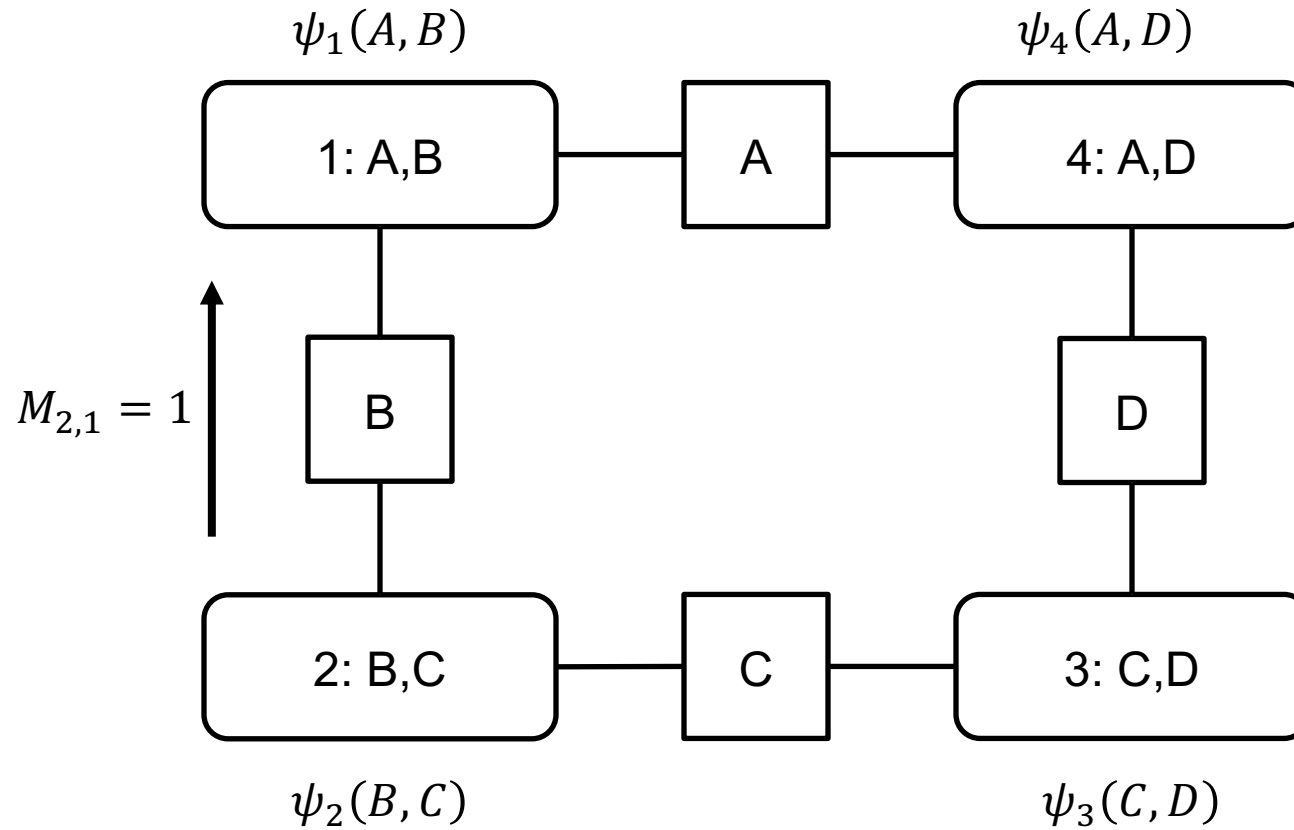
$$\psi_4(A, D) = \phi_4(A, D)$$



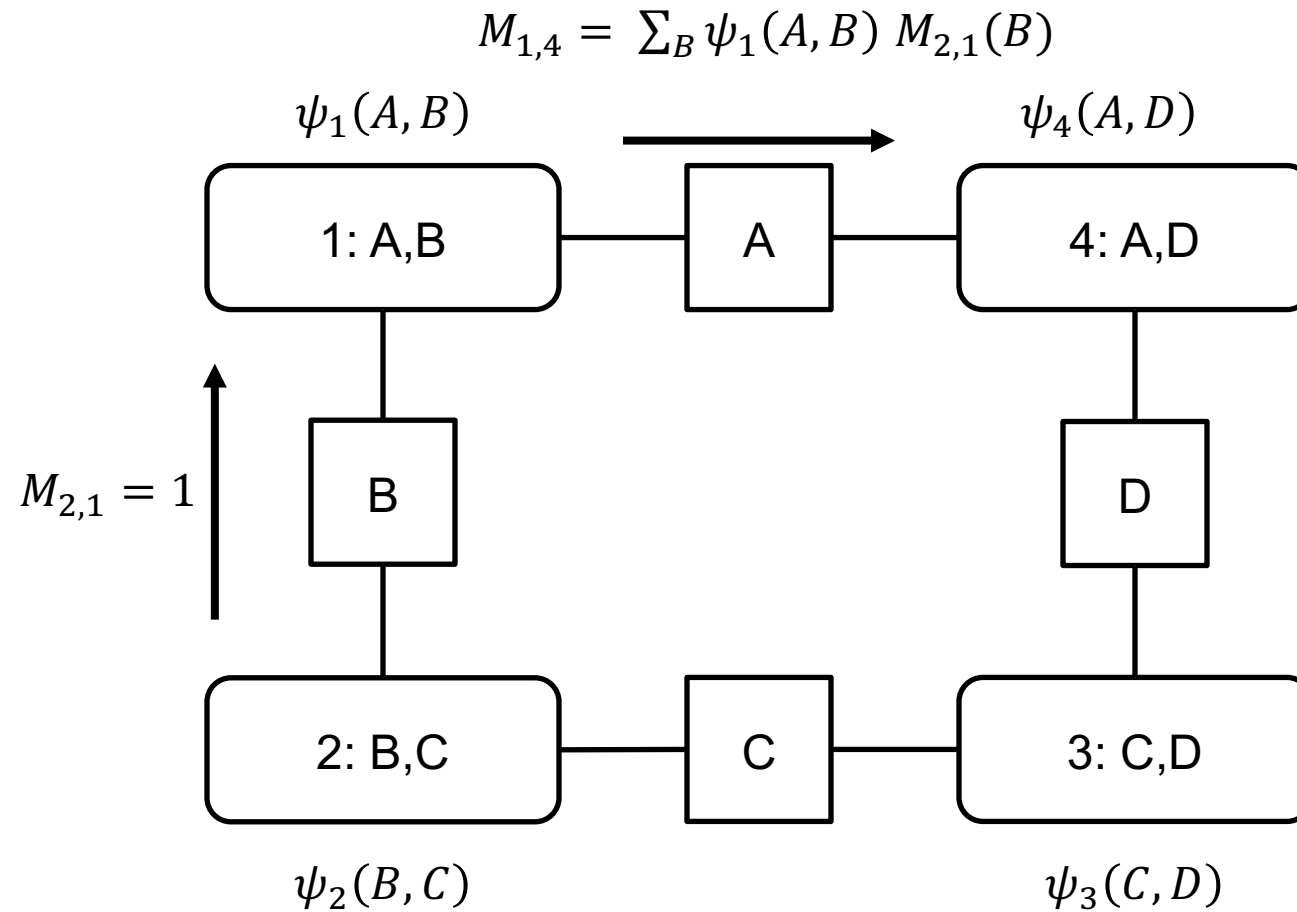
$$\psi_2(B, C) = \phi_2(B, C)$$

$$\psi_3(C, D) = \phi_3(C, D)$$

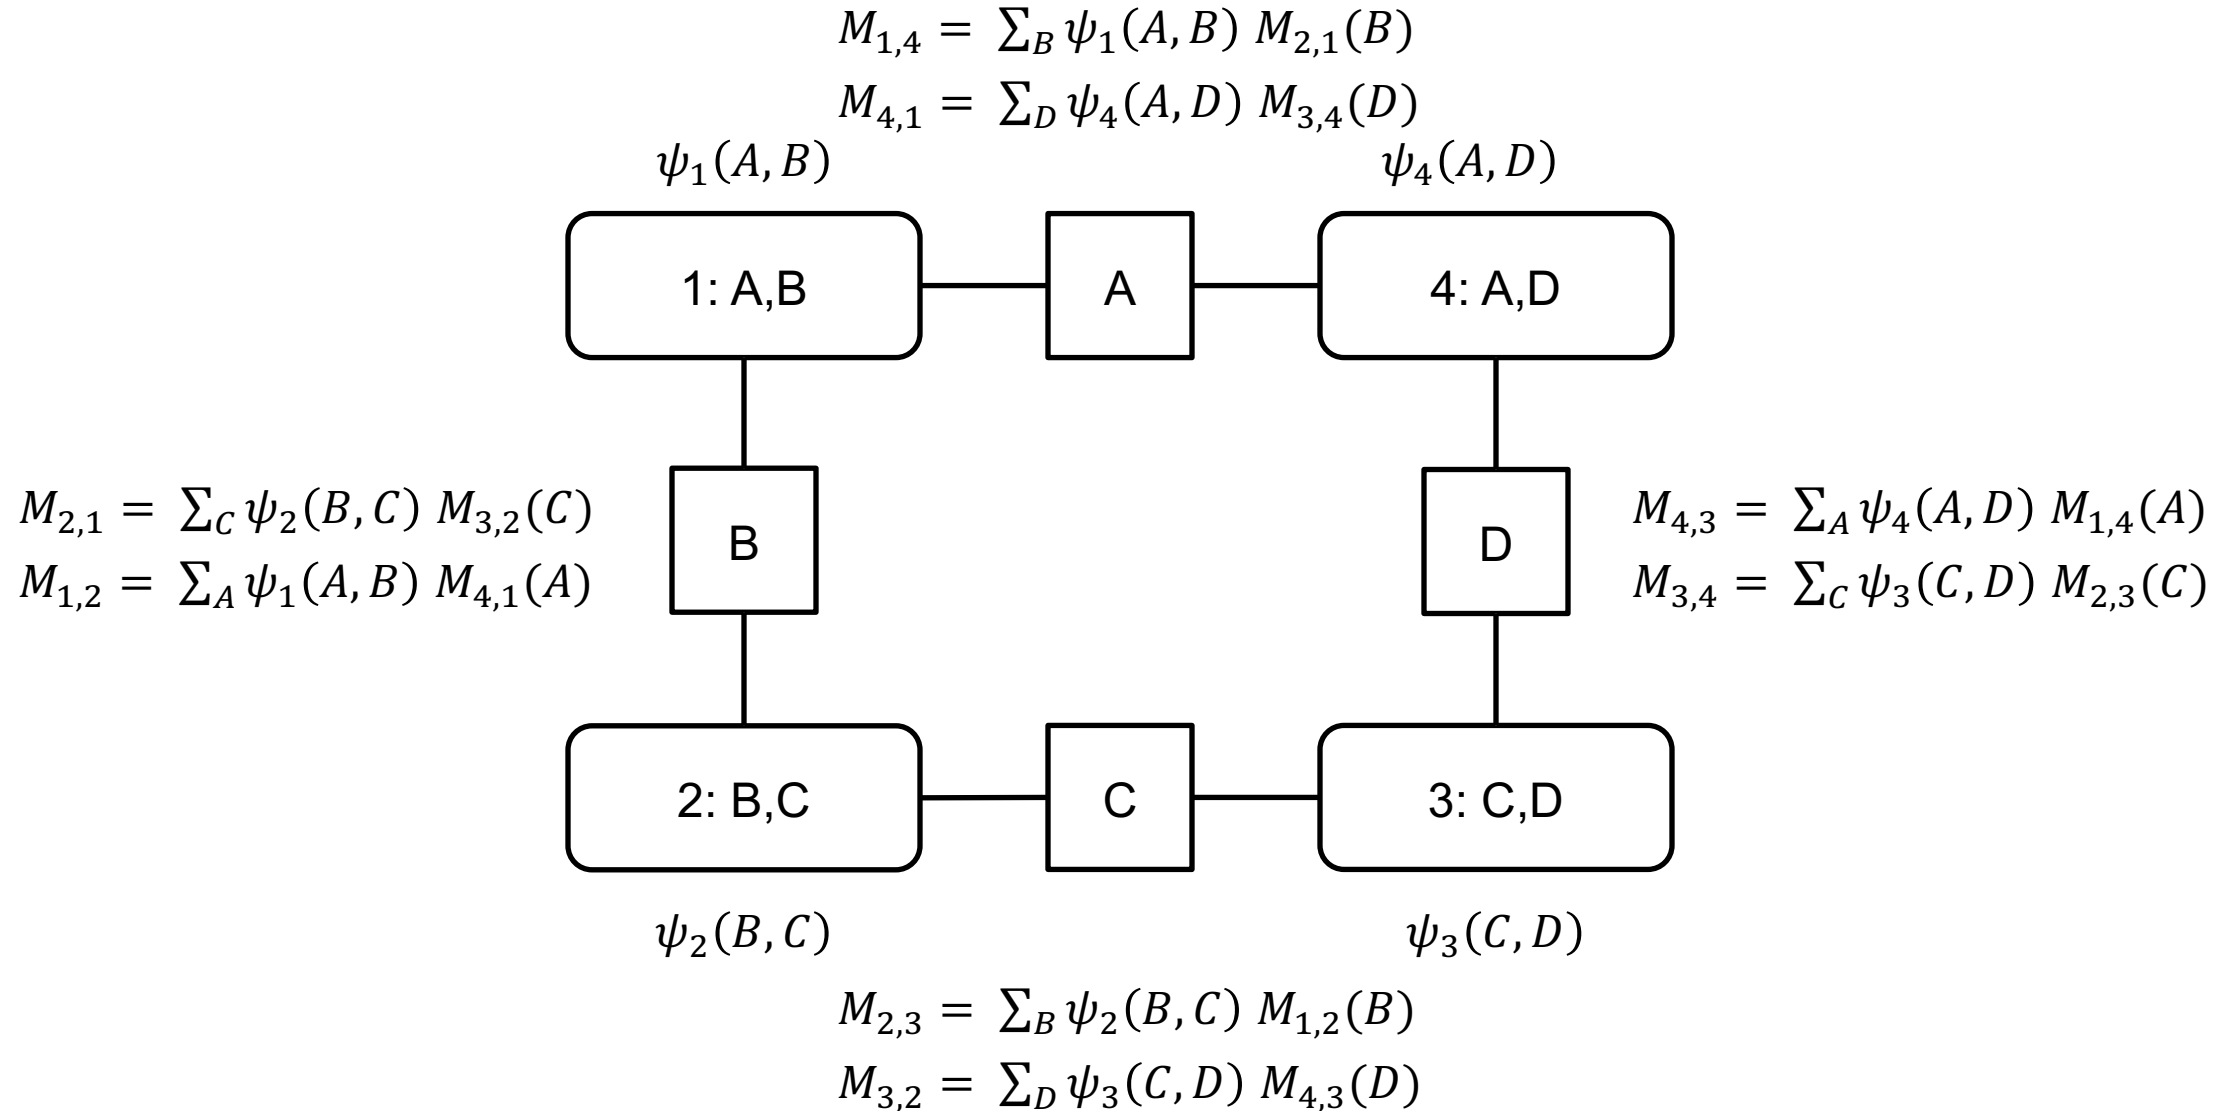
Joingraph Example with Markov Nets



Message Passing with Markov Nets



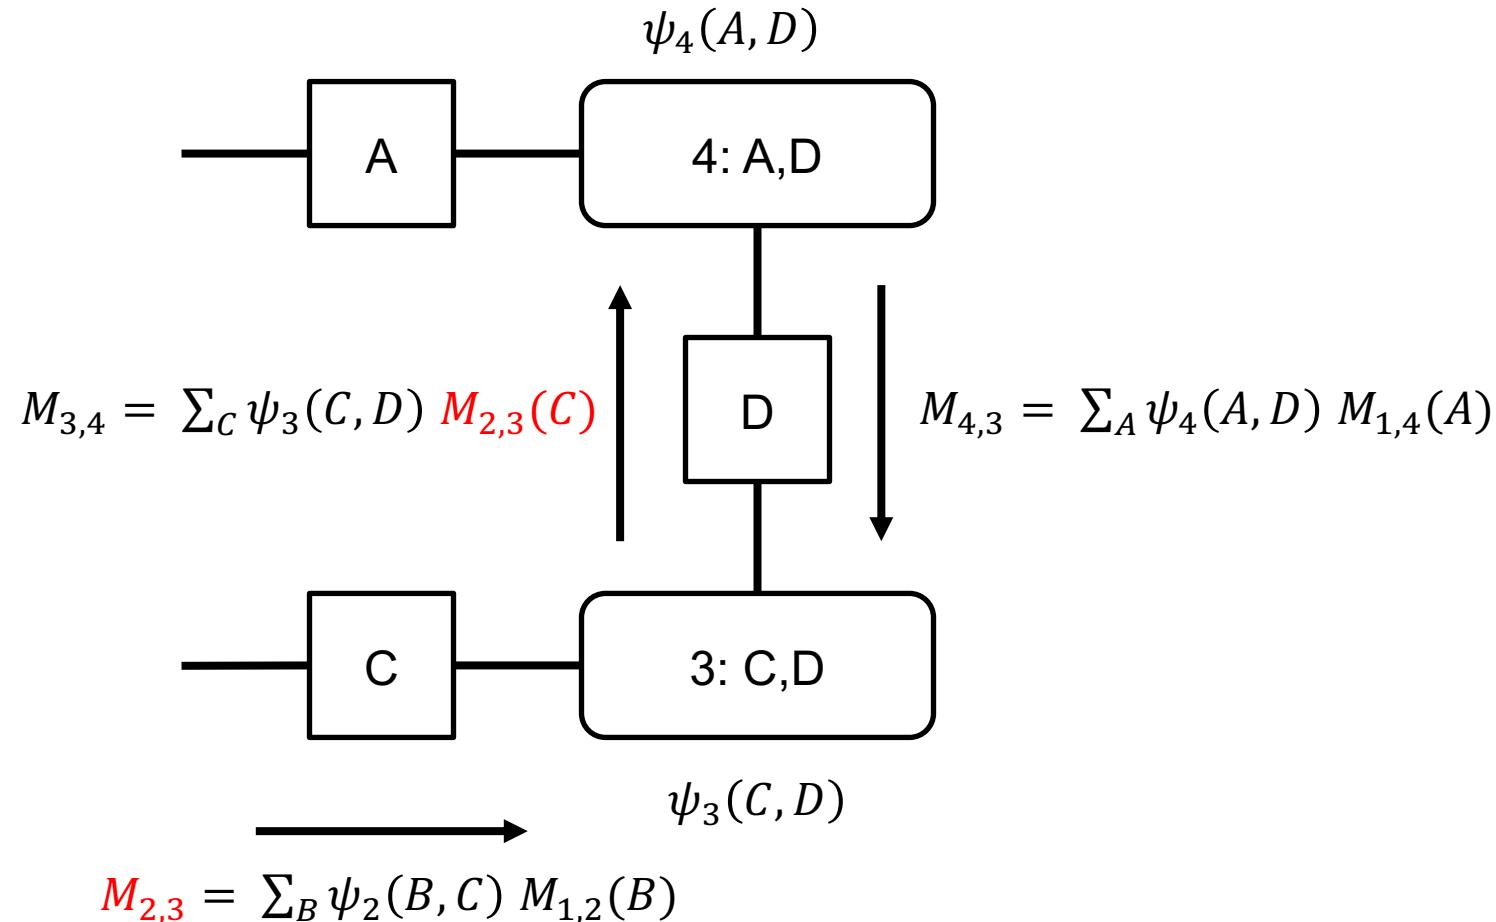
Message Passing with Markov Nets



Message Passing: Avoid Self-Beliefs

- Notice that


- $M_{3,4}$ only considers information received from 2 ($M_{2,3}$)
- Therefore, ignoring information from 4
- This helps to avoid reinforcing self-beliefs



Bethe Graph

- A simple way to generate a valid joingraph
 - For each ϕ_k , make a cluster $\mathcal{C}_k = Vars(\phi_k)$
 - For each variable X_i , create a singleton cluster $\{X_i\}$
 - Edge (\mathcal{C}_k, X_i) if $X_i \in \mathcal{C}_k$


Which factor goes here?
 $\psi_i = 1$





$\phi_1(A, B, C), \phi_2(B, C), \phi_3(B, D), \phi_4(D, E), \phi_5(B, E), \phi_6(B, D), \phi_7(B, D, F)$

Belief Propagation Error

Pyramid network




Loopy belief propagation





Likelihood sampling with 200 cases

Belief Propagation Error

toyQMR network




Loopy belief propagation





Likelihood sampling with 200 cases

Belief Propagation Error

Full scale ICU network



Loopy belief propagation





Likelihood sampling with 200 cases

Belief Propagation Error

QMR-DT network


Posterior
marginals for
three nodes



Exact marginals
(circles) and
error bars


Convergence and Message Schedule

- In their analysis, Murphy, Weiss & Jordan used synchronous (parallel) message passaging
 - However, convergence can be improved with asynchronous approaches
- Some approaches for asynchronous message scheduling
 - Tree reparameterization (TRP): Choose a tree (spanning tree is a good choice) and pass messages. The trees must cover all edges
 - Residual belief propagation (RBP): Pass messages between two clusters whose beliefs over separators disagree the most. Usually organised with a priority queue
- Smoothing messages
 - $$M_{ij} = \lambda \left(\eta \sum_{c_i \setminus s_{ij}} \psi_i \prod_{k \neq j} M_{ki} \right) + (1 - \lambda) M_{ij}^{old}$$




Joingraph Example with Markov Nets


D	A	$\phi_4(D, A)$	A	B	$\phi_1(A, B)$
d	a	100	a	b	100
d	\bar{a}	1	a	\bar{b}	2
\bar{d}	a	1	\bar{a}	b	1
\bar{d}	\bar{a}	100	\bar{a}	\bar{b}	100



C	D	$\phi_3(C, D)$	B	C	$\phi_2(B, C)$
c	d	1	b	c	100
c	\bar{d}	100	b	\bar{c}	1
\bar{c}	d	100	\bar{b}	c	1
\bar{c}	\bar{d}	1	\bar{b}	\bar{c}	100



Convergence and Message Schedule



50 random grids of size 11×11 and $C = 11$ (left) and $C = 13$ (right)

Conclusion

- Belief propagation extends the paradigm of message passing
 - It provides a full spectrum of possibilities from exact to approximate inference
- Interactive joingraph propagation (IJGP) algorithm
 - Can be interpreted as an approach that minimizes the KL divergence between
 - The factorization induced by the network
 - The factorization induced by the joingraph
- IJGP messages convergence
 - Guaranteed in a single iteration if the joingraph is a tree (jointree)
 - Otherwise, convergence is not guaranteed
 - Even if the messages converge, its beliefs may not be necessarily equal the true marginals
 - Although very often in practice they will be close
- Task
 - Read chapter 14 (but 14.8)